

W4111 Introduction to Databases
Fall 2015
Midterm Exam Solutions

Closed Book, 1/2 sheet paper both sides
Duration: 75 minutes

Professor: Eugene Wu
Wednesday, Oct 28, 2015

Your Name: Alice H. Acker

Your UNI: _____

UNI of person to your left: _____

UNI of person to your right: _____

Problem	Points	Score
1	18	18
2	28	28
3	26	26
4	10	10
5	16	16
6	0	3
Write your UNI in the top right right corner of each page.	2	2
TOTAL	100	103

Good Luck!




Figure 1: Cumulative Distribution of the exam scores




Figure 2: Histogram of the exam scores

Note that the solutions discuss the most common case solutions.

5 students recieved extra credit. The fact that the distribution resembles a normal distribution, and the mean and median are the same is a good sign:

	Avg	Median	Stddev	Max	Min
Actual	54	54	14.3	84	16
Estimated, personal	58				
Estimated, course avg	61.8				

For regrade requests

If we added your score wrong, let us know in OH and we will correct it.

If we miscalculated your extra credit, let us know in OH and we will correct it.

If you want a regrade of a subset of the questions, we will *regrade the entire exam* carefully.

1 (18 points) Terms and Definitions

(3 points each) In *at most* two short sentences each, explain the meaning of the following terms as they relate to database management systems. **Draw a box around your final answer.**

1. Total Participation

Used to represent at-least-one relationships between two entities.

2. Transaction

Ensures that a sequence of statements happens all at once or not at all in a single instant in time.

3. Bag/Multi-Set Semantics

Duplicates are allowed in relational tables. SQL uses these semantics, thus necessitating the DISTINCT operator.

Note: has nothing to do with “multiple-sets”. The multi term describes the multiplicity of member elements in a set.

4. Positional-field Notation

Used to unambiguously refer to attributes by their position. Useful when there is ambiguity in column names (common in intermediate query results).

5. User Defined Functions

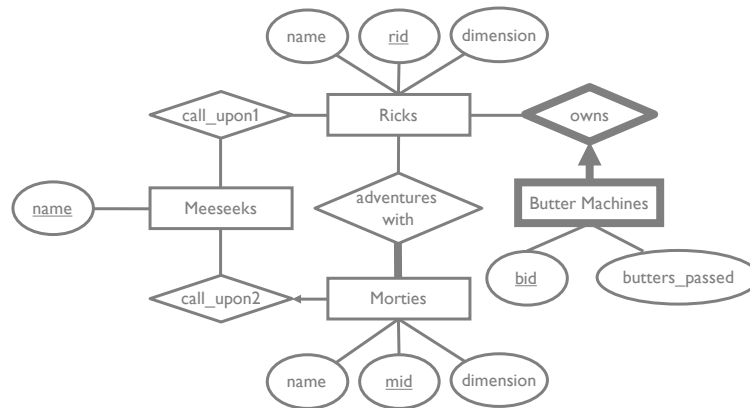
Allows users to write custom programmatic functions that can be used in SQL queries.

6. Physical-data Independence

Protects the application from changes in the physical schema of the database. A core property of declarativity.

2 (28 points) Bulitng a Multi-verse using ER diagrams

Consider the following Entity-Relationship diagram that describes relationships in Rick and Morty's world. There are multiple Ricks and Morties, each from a particular dimension. Ricks go on multi-dimensional adventures with Morties. Ricks also own butter machines whose job is to pass the butter. In addition, Meeseeks are creatures that can be called upon by Ricks and/or Morties to accomplish tasks – in fact, that is their sole purpose. The description above does not describe the multiplicity of the relationships in the diagram, which is your job to understand!



1. (12 points) For the next 4 items, answer them based on the ER diagram above. You will get 3 points for each correct answer, -0.5 points for each incorrect answer, and 0 points for each answer left blank.

- (a) **True / False** There are Morties that grow up and live their lives without going on any adventures with a Rick.

FALSE. There is a total-participation relationship.

- (b) **True / False** A single Meeseek can be called upon by both a Rick from dimension C-137 and a Morty from dimension C-137.

TRUE. The relationship between a Meeseek and a Rick is separate from the relationships between that Meeseek and a Morty.

- (c) **True / False** Ricks have a lot of enemies (including other Ricks) and might be exterminated, thus removing the dead Rick from the Ricks table. In that situation, does this allow the deceased Rick's Butter Machine to run away and live a long, prosperous, non-butter-passing life?

FALSE Sadly, the Butter Machine is a weak entity and cannot be represented without its owner Rick.

- (d) Consider Morty from dimension C-137, let's call him $Morty_{c137}$. $Morty_{c137}$ is feeling overwhelmed about missing math class to go on adventures and has called upon as many Meeseeks as he can (according to the ER diagram). Circle the most general correct answer.
- $Morty_{c137}$ must have called more Meeseeks than some Ricks.
 - $Morty_{c137}$ must have called more Meeseeks than all Ricks.

- iii. $Morty_{c137}$ must have called less Meeseeks than some Ricks.
- iv. $Morty_{c137}$ must have called less Meeseeks than all Ricks.
- v. None of the above.

iii: less Meeseeks than some Ricks. The arrow means $Morty_{c137}$ can call at most 1 Meeseek, whereas Ricks can call as many as they want. Thus $Morty_{c137}$ must call less than *some* Ricks.

2. (16 points) Consider the following SQL schema, which translates of the ER diagram above into tables in the relational model.

```
CREATE TABLE Ricks(
    rid int UNIQUE,
    name text,
    dimension text
);

CREATE TABLE MortiesAdventuresWith(
    mid int NOT NULL,
    name text,
    dimension text,
    rid int REFERENCES Ricks(rid),
    mname_call_upon2 text REFERENCES Meeseeks2(name),
    PRIMARY KEY (mid)
)

CREATE TABLE ButterMachines(
    bid int,
    butters_passed int,
    check (bid is not null),
    rid int REFERENCES Ricks(rid) NOT NULL ON CASCADE delete,
    PRIMARY KEY (bid, rid)
)

CREATE TABLE Meeseeks1(
    name text UNIQUE NOT NULL,
    PRIMARY KEY (name)
);

CREATE TABLE call_upon1(
    rick_id int,
    mname text REFERENCES Meeseeks1(name),
    UNIQUE (rick_id, mname)
);

CREATE TABLE Meeseeks2(
    name text UNIQUE NOT NULL,
    PRIMARY KEY (name)
);
```

(4 points each) List and briefly explain the four most important problems that you find with this schema, in terms of how well it models the E/R diagram above. For each problem, provide a *very* brief example of data that illustrates the problem. Do not base your answer on comparing this schema with other possible schemas. Instead, just compare the schema against the E/R diagram to identify what is not captured from the diagram, etc. in the schema.

You will be graded on the importance of the problems that you identify and the extent that the problems you describe overlap. For example, if you say “table Foo is wrong” as one problem, and “relationship with Foo is a problem because Foo is wrong”, then we will grade the second “problem” lower.

Draw a box around your final answer.

- (a) **Problem 1:**
- (b) **Problem 2:**
- (c) **Problem 3:**
- (d) **Problem 4:**

Any of the following were given points. Superfluous or contradictory text resulted in point reductions. Overlapping problems were given 2 points instead of 4. CASCADE DELETE was a typo addressed during the exam and is not a problem.

- (a) Ricks doesn't have primary key
- (b) ButterMachines.bid should be UNIQUE
- (c) Meeseeks should be a single table
- (d) call_upon1 should have primary key rickid, mname
- (e) MortiesAdventuresWith.rid should be NOT NULL
- (f) Primary key on MortiesAdventuresWith.mid restricts Morties to only one adventure
- (g) In general, total participation cannot be expressed adequately in the relational model
- (h) call_upon1.rick_id should reference Ricks
- (i) ButterMachines may be owned by ≥ 1 Ricks
- (j) Should have CHECK on butterspassed to ensure no negative values

3 (26 points) Save the Multi-verse with SQL

Consider the following intergalactic friendship database, consisting of two relations for `users` and `friends`. The direction of the friendship is explicitly stored, such that the friender A initiated the friendship and the friendee B accepts the friendship. This means B can also initiate a friendship with A, so that B is in two records – one as a friender and one as a friendee. Despite this, we say that users X and Y are friends regardless of which one is the friender and friendee. Users cannot be friends with themselves.

```
CREATE TABLE users(  
    uid int PRIMARY KEY,  
    username text UNIQUE NOT NULL,  
    species text  
);  
  
CREATE TABLE friends(  
    friender int REFERENCES users(uid),  
    friendee int REFERENCES users(uid),  
    when_added date,  
    PRIMARY KEY (friender, friendee)  
);
```

In the following 3 questions, write **a single** SQL statement that satisfies the english statement. You may use SQL features such as the `WITH` clause. Draw a box around your final answer. In the last question, you will fill in a `TRIGGER` template to express the english statement.

1. **(6 points)** Warren-G and Ice-T are Alphabetrians (users of the species = Alphabetrian). Find the usernames of all users that are friendees of both of these Alphabetrians.

- (a) +2 points: if correct query for friends of warren-g
(b) +2 points: if correct query for friends of ice-t
(c) full points: if correct query

```
SELECT u2.username  
FROM users u1, users u2, friends  
WHERE u1.uid = friends.friender and  
      u2.uid = friends.friendee and  
      u1.username = 'Warren-G'  
INTERSECT  
SELECT u2.username  
FROM users u1, users u2, friends  
WHERE u1.uid = friends.friender and  
      u2.uid = friends.friendee and  
      u1.username = 'Ice-T'
```

2. **(6 points)** Alphabetrians are at war with Numbericons (users of the species = Numbericon) and it's time to *crunch the numbers*. We need to list our Alphabetrian brethren and note down the traitors (users that are friends with a Numbericon). Create a two column table containing all Alphabetrians. The first column is the username and the second column is `TRUE` if the user is a traitor, and `NULL` or `FALSE` otherwise. The names of the columns don't matter.

- (a) +1 point: if correct query for alphabetrian users

- (b) +2 points: if correct query for numbericon friends
- (c) +1 points: if used outer join or equivalent query structure (see above)
- (d) full points: if correct
- (e) no deductions if spelled numbericon incorrectly

```

WITH numbericonFriends(uid) AS (
  SELECT bidirfriends.dst
  FROM (
    SELECT friender as src, friendee as dst FROM friends
    UNION
    SELECT friendee as src, friender as dst FROM friends
  ) as bidirfriends,
  users
WHERE bidirfriends.src = users.uid and
      users.species = 'Numbericon'
)
(
  SELECT users.username, numbericonFriends.uid is not null
  FROM users LEFT OUTER JOIN numbericonFriends
  ON users.uid = numbericonFriends.uid
)

```

OR

```

WITH numbericonFriends(uid) AS (
  SELECT bidirfriends.dst
  FROM (
    SELECT friender as src, friendee as dst FROM friends
    UNION
    SELECT friendee as src, friender as dst FROM friends
  ) as bidirfriends,
  users
WHERE bidirfriends.src = users.uid and
      users.species = 'Numbericon'
)
(
  SELECT users.username, users.uid IN (numbericonFriends)
  FROM users
)

```

3. **(8 points)** Rick is wanted by the Inter-dimensional Federation and we're going to use social media to crack down on his rebel group. Find the usernames of all users that are frienders of all of Rick's friends (Rick has username = Rick Sanchez)

- (a) +2 point: if correct query for finding rick's friends (both as friender and friendee)
- (b) +4 points: if correct double negation structure for the "frienders of ALL of Rick's friends".
- (c) full points: if correct
- (d) no deductions if spelled numbericon incorrectly

```

WITH rickfriends(uid) AS (
  -- rick is friendee
  SELECT friends.friender
  FROM friends JOIN users ON users.uid = friends.friendee
  WHERE users.username = 'Rick Sanchez'
  UNION
  -- rick is friender

```



```

        SELECT friends2.friendee
        FROM friends as friends2 JOIN users users2 ON users2.uid = friends2.friender
        WHERE users2.username = 'Rick Sanchez'
    )

    SELECT users.username
    FROM users
    WHERE NOT EXISTS (
        SELECT *
        FROM rickfriends
        WHERE NOT EXISTS (
            SELECT *
            FROM friends
            WHERE users.uid = friends.friender AND
                  rickfriends.uid = friends.friendee
        )
    )
)

```

4. (6 points) Fill in the following logical TRIGGER definition (not postgresql) using PL/SQL to prevent Cromulons from adding themselves (users of the species = Cromulon) to the users table.

```
CREATE TRIGGER nocromulons
```

```
_____ INSERT ON _____
```

```
FOR EACH _____
```

```
BEGIN
```

```
END;
```

- (a) +2 point: if BEFORE INSERT ON USERS
 (b) +2 points: if FOR EACH ROW
 (c) +2: if RETURN NULL; We accepted minor syntax errors (forget END IF;) and statements without the RETURN NEW.

```

CREATE TRIGGER nocromulons
BEFORE INSERT ON users
FOR EACH ROW
BEGIN
    IF NEW.species = 'Cromulon'
        RETURN NULL;
    END IF;
    RETURN NEW;

END;

```

OR

```
CREATE TRIGGER nocromulons
AFTER INSERT ON users
FOR EACH ROW [or STATEMENT]
BEGIN
    DELETE FROM users WHERE species = 'Cromulon';
    RETURN NULL;
END;
```

4 (10 points) Inter-dimensional Relations with Relational Algebra

Consider the database in Problem 3. Write the relational algebra query to find the usernames of users that are frienders with ALL users of species = Alphabeticrian, but NOT frienders with species = Numbericon.

Draw a BOX around your final and intermediate answers – partial credit will be given to clearly labeled intermediate expressions

1. **2 points** got alphabeticrians

$p(\text{alphasfrienderes}(\text{uid} \rightarrow \text{friendee}), \pi_{\text{uid}}(\sigma_{\text{species}=\text{Alphabeticrian}}(\text{users})))$

2. **2 points** Division for “frienders with ALL alphabeticrians” correct

$p(\text{alphas}, \pi_{\text{friendee}, \text{friendee}}(\text{friends}) / \text{alphasfrienderes})$

3. **2 points** Frienders of numericons

$p(\text{numbers}, \pi_{\text{friendee}}(\sigma_{\text{species}=\text{Numbericon}}(\text{friends} \bowtie_{\text{uid}=\text{friendee}} \text{users})))$

4. **1 points** Exclude numericons

$p(\text{tmp}, \text{alphas} - \text{numbers})$

5. **3 points** Get the final user name aka correct answer

$\pi_{\text{username}}(\text{tmp} \bowtie_{\text{uid}=\text{friendee}} \text{users})$

NOTE: The $\bowtie_{\text{uid}=\text{friendee}} \text{users}$ in the final query (5) may be instead run in the definition of *alphas* and *numbers*.

5 (16 points) A =? B

Consider the relational model and two relations with the same schema, R and S. You may assume that no primary keys have been explicitly defined by the database administrator, but otherwise do not make any assumptions about the database.

R(a int, b int)
S(a int, b int)

Each row of the following table shows two queries. In the blank third column of the table write YES if the two queries are equivalent, and NO if they are not equivalent. Two queries are equivalent if they always return exactly the same result relation on all databases with the schema above.

You will receive **4 points** for each correct answer, **-1 points** for each incorrect answer, **0 points** for each answer left blank.

Q1	Q2	Equivalent?
$\pi_{R.a}((R \cup S) - S)$	$\pi_{R.a}(R) - \pi_{R.a}(R \cap S)$	No. Consider the following example: $R = \{(1, 2), (1, 3)\}$, $S = \{(1, 2), (1, 4)\}$
$\pi_{R.b}(\sigma_{R.a=1}(R))$	SELECT R.b FROM R WHERE R.a = 1	No. Q1 returns a set, Q2 returns a multiset (has duplicates)
$R \bowtie S$	$\sigma_{R.a=S.a}(R \times S)$	No. Q1 is a natural join (on a and b). Q2 is a join only on a
SELECT * FROM R WHERE R.a < ALL (SELECT a FROM S)	SELECT R.a, R.b FROM R JOIN S WHERE R.a < S.a	No. Q1 returns R rows less than ALL S.a values. Q2 returns R rows less than ANY S.a value.

6 (Extra credit) Wisdom of the Crowds

The following questions are for extra credit. Each answer should be a single integer number.

1. (1.5 points) Guess to within 1 point what your midterm exam score will be when the exams are handed back (not including the extra credit).

100

2. (1.5 points) Guess to within 1 point what the class average exam score will be when the exams are handed back (not including the extra credit).

54
