Lecture 2
Entity-Relationship Model

Eugene Wu
Fall 2015

HW 0

VM vs virtualenv vs folder vs python shell

Maybe talk about GIT workflows

How to ask for help
https://github.com/w4111/syllabus#help

The TA Situation

HW0: Why the different results?

Result: 69

```
import csv
file = open('iowa-liquor-sample.csv')
file_reader = csv.reader(file)
n = 0
for row in file_reader:
  for el in row:
  if "single malt scotch" in el.lower():
    n += 1
print n
```

Result: 51

```
file = open('iowa-liquor-sample.csv','r')
n = 0
for line in file:
  temp = line.lower()
  if 'single malt scotch' in temp:
    n += 1
print n
```

HW0 Stats

enrolled 79
on waitlist 84

http://eugenewu.net/students.html

Steps for a New Application

Requirements
  what are you going to build?
Conceptual Database Design
  pen-and-pencil description
Logical Design
  formal database schema
Schema Refinement:
  fix potential problems, normalization
Physical Database Design
  use sample of queries to optimize for speed/storage
App/Security Design
  prevent security problems

Steps for a New Application

Requirements
  what are you going to build?
Conceptual Database Design
  pen-and-pencil description
Logical Design
  formal database schema
Schema Refinement:
  fix potential problems, normalization
Physical Database Design
  use sample of queries to optimize for speed/storage
App/Security Design
  prevent security problems

ER Modeling

Database Apps Are Complicated

Typical Fortune 100 Company

- ~10k different information (data) systems
- 90% relational databases (DBMSes)
- Typical database has >100 tables
- Typical table has 50 – 200 attributes

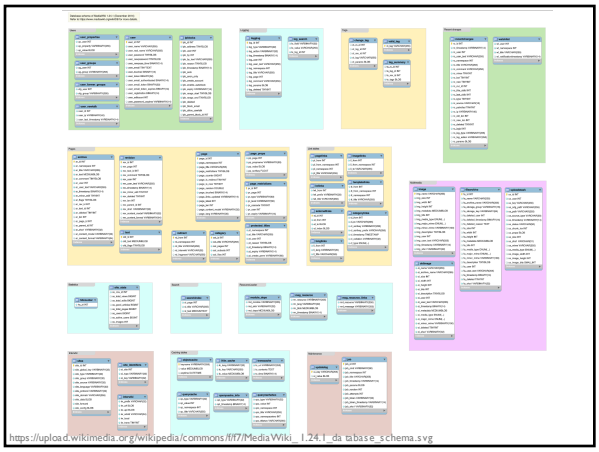

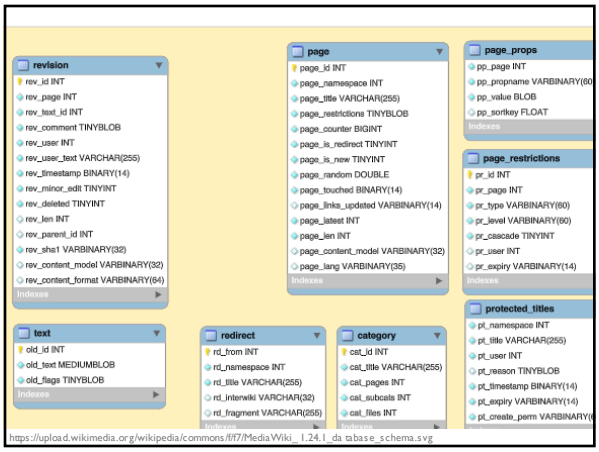

Inconsistencies/Constraint Violations

Huge amount of effort to avoid inconsistencies

DBLP is *the* site for computer science publications

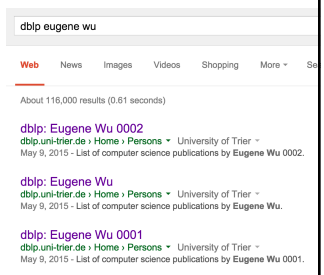

Inconsistencies/Constraint Violations

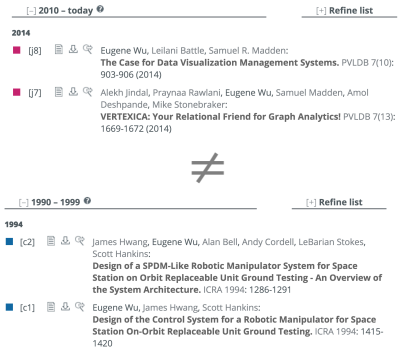

Inconsistencies/Constraint Violations

Giving me eugenewu@gmail would violate constraints

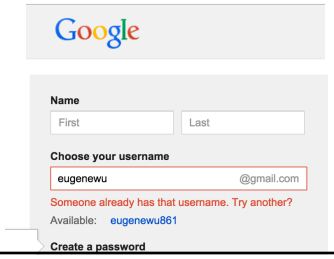

It is Hard to Design Applications

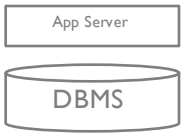

It is Hard to Design Applications

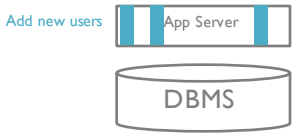

It is Hard to Design Applications

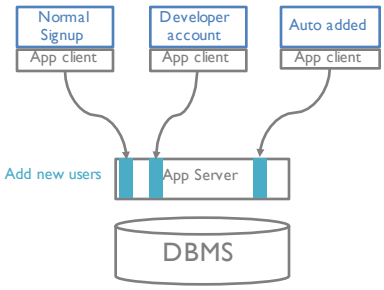

It is Hard to Design Applications

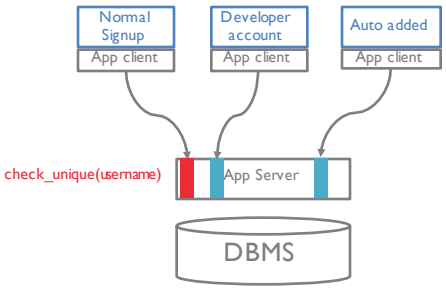

Let's make ~~a webapp~~ $$$

live exercise time

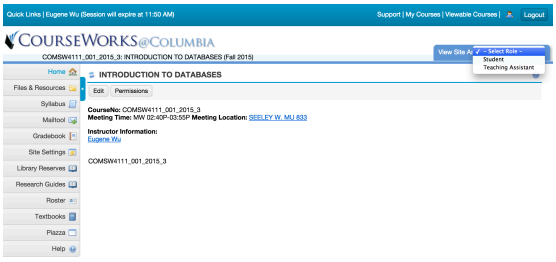

Entity-Relationship Modeling

Entities (objects) to store and their attributes
Relationships between entities and their attrs.
Integrity constraints & business rules
Visually modeled, easy to turn into DB schema

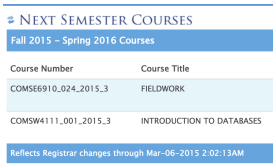
NEXT SEMESTER COURSES

Fall 2015 - Spring 2016 Courses

Course Number	Course Title
COMSE6910_024_2015_3	FIELDWORK
COMSW4111_001_2015_3	INTRODUCTION TO DATABASES

Reflects Registrar changes through Mar-06-2015 2:02:13AM

Courses
- Course Number
- Course Title
- Year
- Semester

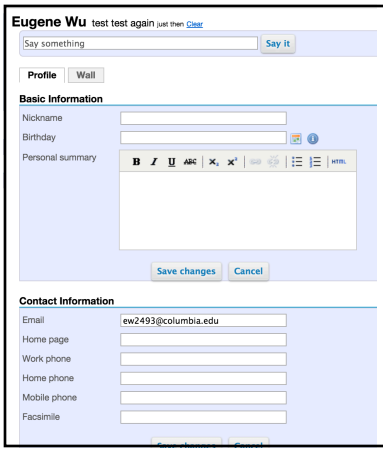

Users
- Nickname
- Name
- Birthday
- Summary
- Email
- …

Basics: Entities

Entity e.g., intro to databases
- real-world object distinguishable from other objects
- described as set of attribute & the values
- (think one record)

Entity Set e.g., all courses
- collection of similar entities
- all entities have same attributes (unless Is-A)
- must have one or more keys
- attributes have domains
- ≈ table

Example: Entity

Keys (cid, uid) are underlined
- Values must be unique
- (think: can use as hashtable key to lookup table)

Basics: Relationships

Relationship: association between 2 or more entities
- e.g., alice is taking Introduction to DBs

Relationship Set: collection of similar relationships
- N-ary relationship set R relates N entity sets $E_1 \ldots E_n$
- Each $r \in R$ involves entities $e_1 \ldots e_n$
- An E_i can be part of diff. relationship sets or diff. roles in same set

Basics: Relationships

Users takes diff roles in same relationships set

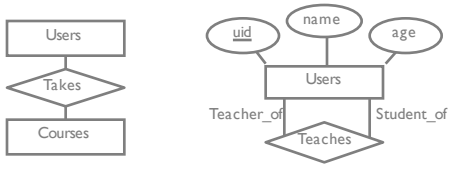

Basics: Relationships

Relationships sets can have descriptive attributes
e.g., the *since* attribute of Instructs

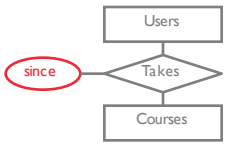

Basics: Ternary Relationships

Connects three entities
N-ary relationships possible too.

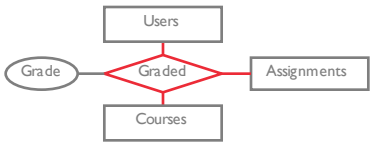

Constraints

Help avoid corruption, inconsistencies

- Key constraints
- Participation constraints
- Weak entities
- Overlap and covering constraints

Key Constraints

Defines cardinality requirements on relationships
Many to many e.g., consider *Takes*
a user can take many courses
a course can have many users that take the course
One to Many e.g., consider *Instructs*
a course has at most one instructor

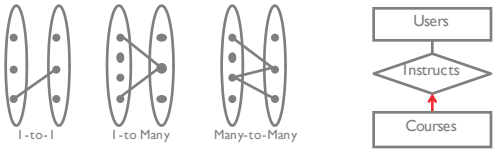

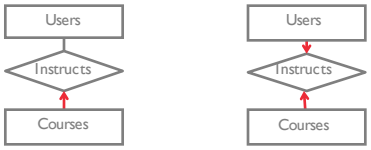

A course has at most one instructor

???

Participation Constraints

Does every course need an instructor?
If yes, it's a participation constraint
e.g., participation of Courses in instructs is *Total*
Otherwise, *partial* participation constraint

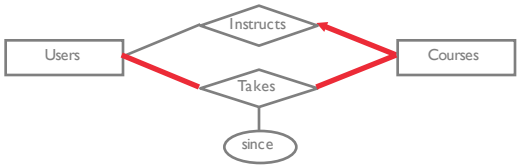

Weak Entities

A *weak entity* can only be uniquely identified by using the primary key of its owner entity

- Owner and weak entity sets must be in one to many relationship set
- Weak entity set must have total participation in this *identifying* relationships set

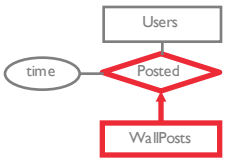

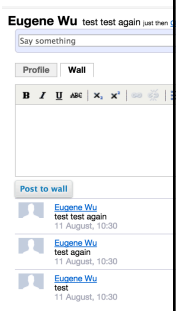

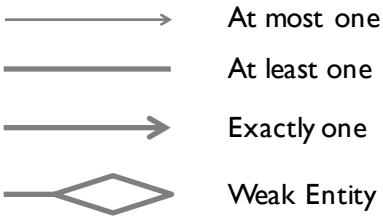

ISA (is a) Hierarchies

Inheritance rules similar to programming languages

- A ISA B → every A also considered a B
- When querying for Bs, must consider As (unlike e.g., C++)

Why use ISA?

- add descriptive attributes specific to a subclass e.g., grade
- identify entities that participate in a relationship

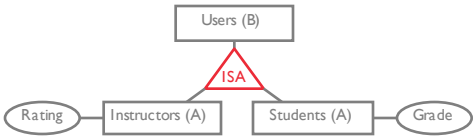

ISA (is a) Hierarchies

Overlap Constraint

- can eugene be an instructor and a student? (allow/disallow)

Covering Constraint

- must every user be an instructor or student? (yes/no)

HOW DO WE EXPRESS THESE IN AN ER DIAGRAM???????

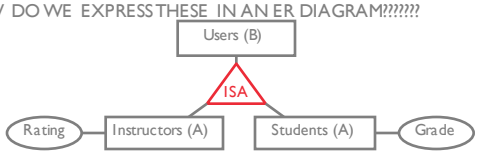

Update on the Course

Enrollment hijinks are over

- 5 Tas! → Increased cap to 120.
- 123 enrolled
- Waitlist auto-maintained (you can see status)

Did you get a github invite?

Project 1 Part 1 out

Aggregation

Relationships between (entities – relationships)

Lets us treat a Relationship Set like an Entity Set

- so it can participate in other relationships

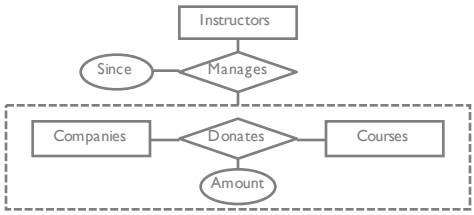

Aggregation vs Ternary Relationships

Why use aggregation?

- Manages and Donates are distinct relationships with own attrs
- Can define constraints on relationship sets
  - e.g., a donation can be managed by at most one instructor

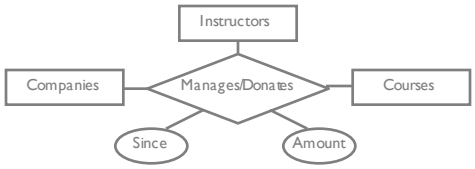

Aggregation vs Ternary Relationships

Constraints apply to all connected entity sets

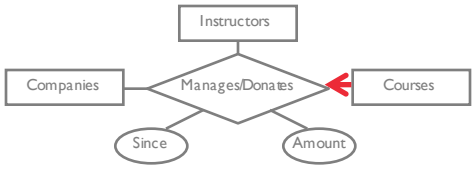

Using the ER Model

Design Choices for a concept

- Entity or Attribute?
- Entity or Relationship?
- Binary or Ternary relationship?
- Aggregation or Ternary relationship?

Entity or Attribute?

Is users.address an attribute of Users or an entity connected to Users by a relationship?

Depends (and may change over time!)

- If a user has >1 addresses, must be an entity
- If an address has attrs (structure), must be entity
  - e.g., want to search for users by city, state, or zip

Entity or Attribute?

A company can't donate multiple amounts (top fig)
Use ternary relationship (bottom fig)

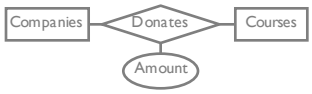

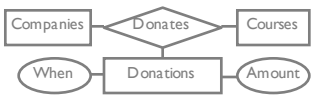

Entity or Relationship?

OK if company donates to courses individually
What if company donates to school for all data-related courses?

- Redundancy of *amount*, need to remember to update every one
- Misleading implies *amount* tied to *each* donation individually

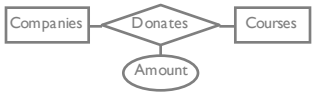

Company	Course	Amount
Amazon	4111	2000
Amazon	4112	2000
Amazon	5111	2000

These amounts are logically the same (redundant)!

Entity or Relationship?

If company donates once to school for data related courses.

Refactor amount into an entity

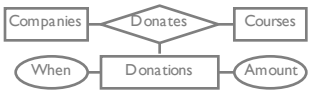

Company	Course	Donation
Amazon	4111	1
Amazon	4112	1
Amazon	5111	1

Donation	When	Amount
1	Today	2000

Binary or Ternary Relationship?

What if assignments have at most one grader?

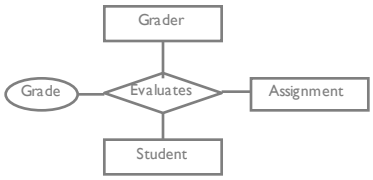

Binary or Ternary Relationship?

What if assignments have at most one grader?

Only one student can complete HW0!

Actually two separate relationships

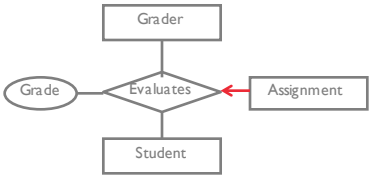

Binary or Ternary Relationship?

Binary relationships allows additional constraints

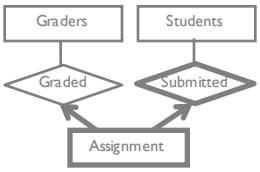

Binary or Ternary Relationship?

Sometimes have true ternary relationship that is defined by all three entities.

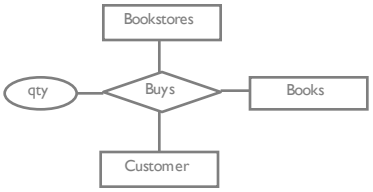

Using ER Modeling

Constraints in ER Modeling

- Many types of data semantics can be captured using ER
- Some constraints not captured (discuss limitations later)

Need further schema refinement

- ER Model is still subjective, need further refinement after translated into relational schema

Summary

Requirements
- what are you going to build?

Conceptual Database Design
- pen-and-pencil description

(Today) ER Modeling

Logical Design
- formal database schema

Schema Refinement:
- fix potential problems, normalization

Physical Database Design
- use sample of queries to optimize for speed/storage

App/Security Design
- prevent security problems

Summary

Conceptual design follows *requirements analysis*

ER model helpful for conceptual design
- constraints are expressive
- matches how we often think about applications

Core constructs
- entity, relationship, attribute
- weak entities, ISA, aggregation

Many variations beyond today's discussion

Summary

ER design is subjective based on usage+needs
- Today we saw multiple ways to model same idea

ER design is not complete/perfect
- Developed in an enterprise-oriented world (ER First)
- Doesn't capture semantics (what does "instructor" *mean?)*
- Doesn't capture e.g., processes/state machines
- How to combine multiple ER models automatically?
- Limitation of imagination when designing webapp
- Open problems!

ER design is a useful way of thought

Next Time

Relational Model: de-facto DBMS standard

Set up for ER diagrams → Relational models

- Maybe talk about abstractions