

L4 Relational Algebra

Eugene Wu
Fall 2015

Overview

Last time, learned about
pre-relational models
an informal introduction to relational model
an introduction to the SQL query language.

Learn about formal relational query languages
Relational Algebra (algebra: perform operations)
Relational Calculus (logic: are statements true?)

Keys to understanding SQL and query processing

Who Cares?

Query Semantics

- Concise *domain specific language (DSL)*
- Rich program analysis
- Helps/enables optimization
- Opens up rich set of topics
 - Materialized views
 - Data lineage/provenance
 - Query by example
 - ...

Reading

Ramakrishnan

Optional

Helpful References

https://en.wikipedia.org/wiki/Relational_algebra

What's a Query Language?

Allows manipulation and **retrieval of data** from a database.

Traditionally: QL != programming language
Doesn't need to be turing complete
Not designed for computation
Supports easy, efficient access to large databases

Recent Years

- Scaling to large datasets is a reality
- Query languages are a powerful way to
 - think about data algorithms scale
 - think about asynchronous/parallel programming

Formal Relational Query Languages

Formal basis for real languages e.g., SQL

Relational Algebra

Operational, used to represent execution plans

Relational Calculus

Logical, describes what data users want
(not operational, fully declarative)

Prelims

Query is a function over **relation instances**

$$Q(R_1, \dots, R_n) = R_{\text{result}}$$

Schemas of input and output relations are *fixed* and well defined by the query Q .

Positional vs Named field notation
 Position easier for formal defs
 one-indexed (not 0-indexed!!!)
 Named more readable
 Both used in SQL

Prelims

Relation (for this lecture)

Instance is a set of tuples

Schema defines field names and types (domains)

Students(sid, name, bday, major, gpa)

How are relations different than generic sets (\mathbb{R})?

Can assume item structure due to schema

Some algebra operations (\times) need to be modified

Will use this later

Relational Algebra Overview

Core 5 operations

PROJECT (π)

SELECT (σ)

UNION (\cup)

SET DIFFERENCE ($-$)

CROSSPRODUCT (\times)

Additional operations

RENAME (ρ)

INTERSECT (\cap)

JOIN (\bowtie)

DIVIDE ($/$)

Instances Used Today: Library

Students, Reservations,
Books

RI

| sid | rid | day |
|-----|-----|-------|
| 1 | 101 | 10/10 |
| 2 | 102 | 11/11 |

Use positional or named
field notation

SI

| sid | name | gpa | age |
|-----|--------|-----|-----|
| 1 | eugene | 4 | 20 |
| 2 | barak | 3 | 21 |
| 3 | trump | 2 | 88 |

Fields in query results are
inherited from input
relations (unless specified)

S2

| sid | name | gpa | age |
|-----|-------|-----|-----|
| 4 | aziz | 3.2 | 21 |
| 2 | barak | 3 | 21 |
| 3 | trump | 2 | 88 |
| 5 | rusty | 3.5 | 21 |

Project

$$\pi_{\langle \text{attr1}, \dots \rangle}(R_{\text{in}}) = R_{\text{result}}$$

Pick out desired attributes (subset of columns)

Schema is subset of input schema in the projection list

$\pi_{\langle a, b, c \rangle}(R_{\text{in}})$ has output schema (a,b,c) with types carried over

Project

S2

| sid | name | gpa | age |
|-----|-------|-----|-----|
| 4 | aziz | 3.2 | 21 |
| 2 | barak | 3 | 21 |
| 3 | trump | 2 | 88 |
| 5 | rusty | 3.5 | 21 |

$$\pi_{\text{name, age}}(S2) =$$

| name | age |
|-------|-----|
| aziz | 21 |
| barak | 21 |
| trump | 88 |
| rusty | 21 |

Project

S2

| sid | name | gpa | age |
|-----|-------|-----|-----|
| 4 | aziz | 3.2 | 21 |
| 2 | barak | 3 | 21 |
| 3 | trump | 2 | 88 |
| 5 | rusty | 3.5 | 21 |

$$\pi_{\text{age}}(S2) =$$

| age |
|-----|
| 21 |
| 88 |

Where did all the rows go?
Real systems typically don't remove duplicates. Why?

Select

$$\sigma_{<p>}(R_{in}) = R_{result}$$

Select subset of rows that satisfy condition p
Won't have duplicates in result. Why?
Result schema same as input

Select

S1

| sid | name | gpa | age |
|-----|--------|-----|-----|
| 1 | eugene | 4 | 20 |
| 2 | barak | 3 | 21 |
| 3 | trump | 2 | 88 |

$$\sigma_{\text{age} < 30}(S1) =$$

| sid | name | gpa | age |
|-----|--------|-----|-----|
| 1 | eugene | 4 | 20 |
| 2 | barak | 3 | 21 |

$$\pi_{\text{name}}(\sigma_{\text{age} < 30}(S1)) =$$

| name |
|--------|
| eugene |
| barak |

Union, Set-Difference

$$R_1 \text{ op } R_2 = R_{result}$$

R_1, R_2 must be *union-compatible*

Same number of fields

Field i in each schema have same type

Result Schema borrowed from first arg (R_1)

$R_1(\text{big int, poppa int}) \cup R_2(\text{thug int, life int}) = ?$

Union, Intersect, Set-Difference

S1

| sid | name | gpa | age |
|-----|--------|-----|-----|
| 1 | eugene | 4 | 20 |
| 2 | barak | 3 | 21 |
| 3 | trump | 2 | 88 |

S2

| sid | name | gpa | age |
|-----|-------|-----|-----|
| 4 | aziz | 3.2 | 21 |
| 2 | barak | 3 | 21 |
| 3 | trump | 2 | 88 |
| 5 | rusty | 3.5 | 21 |

$$S1 \cup S2 =$$

| sid | name | gpa | age |
|-----|--------|-----|-----|
| 1 | eugene | 4 | 20 |
| 4 | aziz | 3.2 | 21 |
| 2 | barak | 3 | 21 |
| 3 | trump | 2 | 88 |
| 5 | rusty | 3.5 | 21 |

Union, Intersect, Set-Difference

S1

| sid | name | gpa | age |
|-----|--------|-----|-----|
| 1 | eugene | 4 | 20 |
| 2 | barak | 3 | 21 |
| 3 | trump | 2 | 88 |

S2

| sid | name | gpa | age |
|-----|-------|-----|-----|
| 4 | aziz | 3.2 | 21 |
| 2 | barak | 3 | 21 |
| 3 | trump | 2 | 88 |
| 5 | rusty | 3.5 | 21 |

$$S1 - S2 =$$

| sid | name | gpa | age |
|-----|--------|-----|-----|
| 1 | eugene | 4 | 20 |

Note on Set Difference & Performance

Notice that most operators are monotonic
 increasing size of inputs \rightarrow outputs grow
 if $S1 \supseteq S2 \rightarrow Q(S1, T) \supseteq Q(S2, T)$
 can compute *incrementally*

Set Difference is *not monotonic*

e.g., $T - S1$
 if $S1 \supseteq S2 \rightarrow T - S1 \not\subseteq T - S2$

Set difference is *blocking*:

For $T - S$, must wait for all S tuples before any results

Cross-Product

$$R_1(a_1, \dots, a_n) \times R_2(a_{n+1}, \dots, a_m) = R_{\text{result}}(a_1, \dots, a_m)$$

Each row of $S1$ paired with each row of $R1$

Result schema concatenates $S1$ and $R1$'s fields, inherit if possible

Conflict: $S1$ and $R1$ have *sid* field

Different than mathematical "X" by flattening results:

math $A \times B = \{ (a, b) \mid a \in A \wedge b \in B \}$

e.g., $\{1, 2\} \times \{3, 4\} = \{ (1, 3), (1, 4), (2, 3), (2, 4) \}$

what is $\{1, 2\} \times \{3, 4\} \times \{5, 6\}$?

Cross-Product

| S1 | | | |
|-----|--------|-----|-----|
| sid | name | gpa | age |
| 1 | eugene | 4 | 20 |
| 2 | barak | 3 | 21 |
| 3 | trump | 2 | 88 |

| R1 | | |
|-----|-----|-------|
| sid | rid | day |
| 1 | 101 | 10/10 |
| 2 | 102 | 11/11 |

$S1 \times R1 =$

| (sid) | name | gpa | age | (sid) | rid | day |
|-------|--------|-----|-----|-------|-----|-------|
| 1 | eugene | 4 | 20 | 1 | 101 | 10/10 |
| 2 | barak | 3 | 21 | 1 | 101 | 10/10 |
| 3 | trump | 2 | 88 | 1 | 101 | 10/10 |
| 1 | eugene | 4 | 20 | 2 | 102 | 11/11 |
| 2 | barak | 3 | 21 | 2 | 102 | 11/11 |
| 3 | trump | 2 | 88 | 2 | 102 | 11/11 |

Rename

$p(<\text{new_name}>(<\text{mappings}>), Q)$

Explicitly defines/changes field names of schema

$p(C(I \rightarrow \text{sid1}, 5 \rightarrow \text{sid2}), S1 \times R1)$

$C =$

| sid1 | name | gpa | age | sid2 | rid | day |
|------|--------|-----|-----|------|-----|-------|
| 1 | eugene | 4 | 20 | 1 | 101 | 10/10 |
| 2 | barak | 3 | 21 | 1 | 101 | 10/10 |
| 3 | trump | 2 | 88 | 1 | 101 | 10/10 |
| 1 | eugene | 4 | 20 | 2 | 102 | 11/11 |
| 2 | barak | 3 | 21 | 2 | 102 | 11/11 |
| 3 | trump | 2 | 88 | 2 | 102 | 11/11 |

Compound/Convenience Operators

Convenience operations

INTERSECT (\cap)

JOIN (\bowtie)

DIVIDE ($/$)

Intersect

$$R_1 \cap R_2 = R_{\text{result}}$$

R_1, R_2 must be *union-compatible*

Intersect

| S1 | | | | S2 | | | |
|-----|--------|-----|-----|-----|-------|-----|-----|
| sid | name | gpa | age | sid | name | gpa | age |
| 1 | eugene | 4 | 20 | 4 | aziz | 3.2 | 21 |
| 2 | barak | 3 | 21 | 2 | barak | 3 | 21 |
| 3 | trump | 2 | 88 | 3 | trump | 2 | 88 |
| | | | | 5 | rusty | 3.5 | 21 |

$$S1 \cap S2 =$$

| sid | name | gpa | age |
|-----|-------|-----|-----|
| 2 | barak | 3 | 21 |
| 3 | trump | 2 | 88 |

Intersect

$$R_1 \cap R_2 = R_{\text{result}}$$

R_1, R_2 must be *union-compatible*

Can we express using core operators?

$$S \cap T = ?$$

Intersect

$$R_1 \cap R_2 = R_{\text{result}}$$

R_1, R_2 must be *union-compatible*

Can we express using core operators?

$$S \cap T = S - ? \quad (\text{think venn diagram})$$

Intersect

$$R_1 \cap R_2 = R_{\text{result}}$$

R_1, R_2 must be *union-compatible*

Can we express using core operators?

$$S \cap T = S - (S - T)$$

theta (θ) Join

$$R \bowtie_c S = \sigma_c(R \times S)$$

Most general form

Result schema same as cross product

Often *far* more efficient to compute than cross product

Commutative

$$(A \bowtie_c B) \bowtie_c C = A \bowtie_c (B \bowtie_c C)$$

theta (θ) Join

| S1 | | | | R1 | | |
|-----|--------|-----|-----|-----|-----|-------|
| sid | name | gpa | age | sid | rid | day |
| 1 | eugene | 4 | 20 | 1 | 101 | 10/10 |
| 2 | barak | 3 | 21 | 2 | 102 | 11/11 |
| 3 | trump | 2 | 88 | | | |

$$S1 \bowtie_{S1.sid \leq R1.sid} R1 =$$

| (sid) | name | gpa | age | (sid) | rid | day |
|-------|--------|-----|-----|-------|-----|-------|
| 1 | eugene | 4 | 20 | 1 | 101 | 10/10 |
| 1 | eugene | 4 | 20 | 2 | 102 | 11/11 |
| 2 | barak | 3 | 21 | 2 | 102 | 11/11 |

Equi-Join

$$R \bowtie_{\text{attr}} S = R \bowtie_{R.\text{attr} = S.\text{attr}} S$$

Special case where the condition is attribute equality
Result schema only keeps one copy of equality fields

Natural Join ($R \bowtie S$):

Equijoin on *all* shared fields

Equi-Join

| S1 | | | | R1 | | |
|-----|--------|-----|-----|-----|-----|-------|
| sid | name | gpa | age | sid | rid | day |
| 1 | eugene | 4 | 20 | 1 | 101 | 10/10 |
| 2 | barak | 3 | 21 | 2 | 102 | 11/11 |
| 3 | trump | 2 | 88 | | | |

| $S1 \bowtie_{\text{sid}} R1 =$ | | | | | | |
|--------------------------------|--------|-----|-----|-----|-------|--|
| sid | name | gpa | age | rid | day | |
| 1 | eugene | 4 | 20 | 101 | 10/10 | |
| 2 | barak | 3 | 21 | 102 | 11/11 | |

Division

Let us have relations $A(x,y), B(y)$

$$A/B = \{ \langle x \rangle \mid \exists \langle x,y \rangle \in A \ \forall \langle y \rangle \in B \}$$

Good to ponder, not supported in most systems (why?)

Find all students that have reserved all books

$A/B = \text{all } x \text{ (students) s.t. for every } y \text{ (reservation), } \langle x,y \rangle \in A$

Generalization

y can be a list of fields in B

$x \cup y$ is fields in A

Examples

| A | | R1 | | R2 | | R3 | |
|-----|-----|-----|--|-----|--|-----|--|
| sid | rid | rid | | rid | | rid | |
| 1 | 1 | 2 | | 2 | | 1 | |
| 1 | 2 | | | 4 | | 2 | |
| 1 | 3 | | | | | 4 | |
| 1 | 4 | | | | | | |
| 2 | 1 | | | | | | |
| 2 | 2 | | | | | | |
| 3 | 2 | | | | | | |
| 4 | 2 | | | | | | |
| 4 | 4 | | | | | | |

A/R1

A/R2

A/R3

Is A/B a Fundamental Operation?

No. Shorthand like Joins

joins so common, it's natively supported

Hint: Find all x s not 'disqualified' by some y in B .

x value is *disqualified* if by attaching y value from B , we obtain an xy tuple that is not in A .

| A | | B | |
|-----|-----|-----|--|
| sid | rid | rid | |
| 1 | 1 | 2 | |
| 1 | 2 | 4 | |
| 1 | 3 | | |
| 1 | 4 | | |
| 2 | 1 | | |
| 2 | 2 | | |
| 3 | 2 | | |
| 4 | 2 | | |
| 4 | 4 | | |

Disqualified =
 $A/B =$

| A | | B | | $\pi_x(A) \times B$ | |
|-----|-----|-----|--|---------------------|-----|
| sid | rid | rid | | sid | rid |
| 1 | 1 | 2 | | 1 | 2 |
| 1 | 2 | 4 | | 1 | 4 |
| 1 | 3 | | | 2 | 2 |
| 1 | 4 | | | 2 | 4 |
| 2 | 1 | | | 3 | 2 |
| 2 | 2 | | | 3 | 4 |
| 3 | 2 | | | 4 | 2 |
| 4 | 2 | | | 4 | 4 |
| 4 | 4 | | | | |

Disqualified = $(\pi_x(A) \times B)$
A/B =

| A | | B | | $\pi_x(A) \times B$ | | $\pi_x(A) \times B - A$ | |
|-----|-----|-----|--|---------------------|-----|-------------------------|-----|
| sid | rid | rid | | sid | rid | sid | rid |
| 1 | 1 | 2 | | 1 | 2 | 2 | 4 |
| 1 | 2 | 4 | | 1 | 4 | 3 | 4 |
| 1 | 3 | | | 2 | 2 | | |
| 1 | 4 | | | 2 | 4 | | |
| 2 | 1 | | | 3 | 2 | | |
| 2 | 2 | | | 3 | 4 | | |
| 3 | 2 | | | 4 | 2 | | |
| 4 | 2 | | | 4 | 4 | | |
| 4 | 4 | | | | | | |

Disqualified = $((\pi_x(A) \times B) - A)$
A/B =

| A | | B | | $\pi_x(A) \times B$ | | $\pi_x(A) \times B - A$ | |
|-----|-----|-----|--|---------------------|-----|-------------------------|-----|
| sid | rid | rid | | sid | rid | sid | rid |
| 1 | 1 | 2 | | 1 | 2 | 2 | 4 |
| 1 | 2 | 4 | | 1 | 4 | 3 | 4 |
| 1 | 3 | | | 2 | 2 | | |
| 1 | 4 | | | 2 | 4 | | |
| 2 | 1 | | | 3 | 2 | | |
| 2 | 2 | | | 3 | 4 | | |
| 2 | 2 | | | 4 | 2 | | |
| 3 | 2 | | | 4 | 4 | | |
| 4 | 2 | | | | | | |
| 4 | 4 | | | | | | |

Disqualified = $\pi_x((\pi_x(A) \times B) - A)$
A/B = $\pi_x(A) - \text{Disqualified}$

Names of students that reserved book 2

$$\pi_{\text{name}}(\sigma_{\text{rid}=2} (R1) \bowtie S1)$$

Equivalent Queries

$$\begin{aligned} & p(\text{Tmp1}, \sigma_{\text{rid}=2} (R1)) \\ & p(\text{Tmp2}, \text{Tmp1} \bowtie S1) \\ & \pi_{\text{name}}(\text{Tmp2}) \end{aligned}$$

$$\pi_{\text{name}}(\sigma_{\text{rid}=2} (R1 \bowtie S1))$$

Names of students that reserved db books

Book(bid, type) Reserve(sid, bid) Student(sid)

Need to join with Instructors to access role

$$\pi_{\text{name}}(\sigma_{\text{type}='db'} (\text{Book}) \bowtie \text{Reserve} \bowtie \text{Student})$$

More efficient query

$$\pi_{\text{name}}(\pi_{\text{sid}}((\pi_{\text{cid}}(\sigma_{\text{type}='db'} (\text{Book})) \bowtie \text{Reserve}) \bowtie \text{Student}))$$

Query optimizer can find the more efficient query!

Students that reserved DB or HCI book

- Find all DB or HCI books
- Find students that reserved one of those books

$$\begin{aligned} & p(\text{tmp}, (\sigma_{\text{type}='DB' \vee \text{type}='HCI'} (\text{Book}))) \\ & \pi_{\text{name}}(\text{tmp} \bowtie \text{Reserve} \bowtie \text{Student}) \end{aligned}$$

Alternatives

define tmp using UNION (how?)

what if we replaced \vee with \wedge in the query?

Students that reserved a DB and HCI book

Does previous approach work?

$$p(tmp, (\sigma_{type='DB' \wedge type='HCI'}(Book)))$$

$$\pi_{name}(tmp \bowtie Reserve \bowtie Student)$$

NO

Students that reserved a DB and HCI book

Does previous approach work?

1. Find students that reserved DB books
2. Find students that reserved HCI books
3. Intersection

$$p(tmpDB, \pi_{sid}(\sigma_{type='DB'}(Book) \bowtie Reserve))$$

$$p(tmpHCI, \pi_{sid}(\sigma_{type='HCI'}(Book) \bowtie Reserve))$$

$$\pi_{name}((tmpDB \cap tmpHCI) \bowtie Student)$$

Is the intersection always allowed?

What if it projected book name?

Students that reserved all books

Use division

Be careful with schemas of inputs to / !

$$p(tmp, (\pi_{sid,bid} Reserves) / (\pi_{bid} Books))$$

$$\pi_{name}(tmp \bowtie Student)$$

What if want students that reserved all horror books?

$$p(tmp, (\pi_{sid,bid} Reserves) / (\pi_{bid}(\sigma_{type='horror'}(Book))))$$

Summary

Relational Algebra (RA) operators

Operators are closed (inputs & outputs are relations)

Multiple Relational Algebra queries can have same semantics (always return same results)

Forms basis for optimizations

Next Time

Relational Calculus