

## L5 SQL SQL SQL SQL SQL SQL SQL

Eugene Wu  
Fall 2015

## Didn't Lecture 3 Go Over SQL?

Two sublanguages

**DDL** Data Definition Language  
define and modify schema (physical, logical, view)  
CREATE TABLE, Integrity Constraints

**DML** Data Manipulation Language  
get and modify data  
simple SELECT, INSERT, DELETE  
*human-readable* language

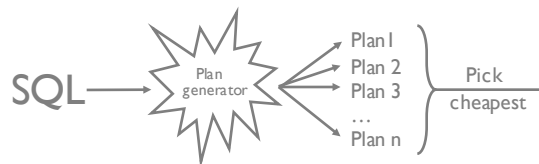
## Gritty Details

**DDL**  
NULL, Views

**DML**  
Basics, SQL Clauses, Expressions, Joins, Nested  
Queries, Aggregation, With, Triggers

## Didn't Lecture 3 Go Over SQL?

DBMS makes it run efficiently  
Key: precise query semantics  
Reorder/modify queries while answers stay same  
DBMS estimates costs for different evaluation plans



## Didn't Lecture 3 Go Over SQL?

More expressive power than Rel Alg  
can be described by extensions of algebra  
**One key difference:** multisets rather than sets  
i.e. # duplicates in a table carefully accounted for

Most widely used *query language*, not just relational  
query language

## Today's Database

Sailors

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Boats

bid	name	color
101	Legacy	red
102	Melon	blue
103	Mars	red

Reserves

sid	bid	day
1	102	9/12
2	102	9/13
2	103	9/14

Is Reserves table correct?

## Today's Database

Sailors

<u>sid</u>	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Boats

<u>bid</u>	name	color
101	Legacy	red
102	Melon	blue
103	Mars	red

Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
1	102	9/12
2	102	9/13
2	103	9/14

Is Reserves table correct?  
Day should be part of key

## Follow along at home!

<http://w4111db1.cloudapp.net:8000/>

or  
psql -U demo -h w4111db1.cloudapp.net demo  
password: demo

## <30 year old sailors

```
SELECT *
FROM Sailors
WHERE age < 30
```

<u>sid</u>	name	rating	age
1	Eugene	7	22
3	Ken	8	27

```
SELECT name, age
FROM Sailors
WHERE age < 30
```

name	age
Eugene	22
Ken	27

## <30 year old sailors

```
SELECT *
FROM Sailors
WHERE age < 30
```

$\sigma_{\text{age} < 30}(\text{Sailors})$

```
SELECT name, age
FROM Sailors
WHERE age < 30
```

$\pi_{\text{name, age}}(\sigma_{\text{age} < 30}(\text{Sailors}))$

## Multiple Relations

```
SELECT S.name
FROM Sailors AS S, Reserves AS R
WHERE S.sid = R.sid AND R.bid = 102
```

Sailors

<u>sid</u>	<u>name</u>	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
1	102	9/12
2	102	9/13
2	103	9/14

## Multiple Relations

```
SELECT S.name
FROM Sailors AS S, Reserves AS R
WHERE S.sid = R.sid AND R.bid = 102
```

$\pi_{\text{name}}(\sigma_{\text{bid}=2}(\text{Sailors} \bowtie_{\text{sid}} \text{Reserves}))$

## Structure of a SQL Query

### DISTINCT

Optional, answer should not have duplicates  
Default: duplicates not removed (multiset)

### target-list

List of expressions over attrs of  
tables in relation-list

```
SELECT [DISTINCT] target-list
FROM relation-list
WHERE qualification
```

### relation-list

List of relation names  
Can define range-variable "AS X"

### qualification

Boolean expressions  
Combined w/ AND, OR, NOT  
attr<sub>1</sub> op const  
attr<sub>1</sub> op attr<sub>2</sub>  
op is =, <, >, !=, etc


## Semantics

```
SELECT [DISTINCT] target-list
FROM relation-list
WHERE qualification
```

FROM compute cross product of relations  
WHERE remove tuples that fail qualifications  
SELECT remove fields not in target-list  
DISTINCT remove duplicate rows

## Conceptual Query Evaluation

```
SELECT [DISTINCT] target-list
FROM relation-list
WHERE qualification
GROUP BY grouping-list
HAVING group-qualification
```



Not how actually executed! Above is likely very slow

## DISTINCT (vol.I)

### Reserves

sid	bid	day
1	102	9/12
2	102	9/13
2	103	9/14

```
SELECT bid
FROM Reserves
```

bid
102
102
103

```
SELECT DISTINCT bid
FROM Reserves
```

bid
102
103

## Sailors that reserved 1+ boats

```
SELECT S.sid
FROM Sailors AS S, Reserves AS R
WHERE S.sid = R.sid
```

Would DISTINCT change anything in this query?  
What if SELECT clause was SELECT S.name?

## Range Variables

### Disambiguate relations

same table used multiple times (self join)

```
SELECT sid
FROM Sailors, Sailors
WHERE age > age
```

```
SELECT S1.sid
FROM Sailors AS S1, Sailors AS S2
WHERE S1.age > S2.age
```

## Range Variables

Disambiguate relations

same table used multiple times (self join)

```
SELECT sid
FROM Sailsors, Sailsors
WHERE age > age
```

```
SELECT S1.name, S1.age, S2.name, S2.age
FROM Sailsors AS S1, Sailsors AS S2
WHERE S1.age > S2.age
```

## Expressions (Math)

```
SELECT S.age, S.age - 5 AS age2, 2*S.age AS age3
FROM Sailsors AS S
WHERE S.name = 'eugene'
```

```
SELECT S1.name AS name1, S2.name AS name2
FROM Sailsors AS S1, Sailsors AS S2
WHERE S1.rating*2 = S2.rating - 1
```

## Expressions (Strings)

```
SELECT S.name
FROM Sailsors AS S
WHERE S.name LIKE 'e_%'
```

'\_' any one character (• in regex)

'%' 0 or more characters of any kind (•\* in regex)

Most DBMSes have rich string manipulation support e.g., regex

PostgreSQL documentation  
<http://www.postgresql.org/docs/9.1/static/functions-string.html>

## Expressions (Date/Time)

```
SELECT R.sid
FROM Reserves AS R
WHERE now() - R.date < interval '1 day'
```

TIMESTAMP, DATE, TIME types

now() returns timestamp at start of transaction

DBMSes provide rich time manipulation support  
 exact support may vary by vendor

Postgresql Documentation  
<http://www.postgresql.org/docs/9.1/static/functions-datetime.html>

## Expressions

Constant	1
Col reference	Sailors.name
Arithmetic	Sailors.sid * 10
Unary operators	NOT, EXISTS
Binary operators	AND, OR, IN
Function calls	abs(), sqrt(), ...
Casting	1.7::int, '10-12-2015'::date

sid of Sailors that reserved red or blue boat

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE B.bid = R.bid AND
      (B.color = 'red' OR B.color = 'blue')
```

OR

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE B.bid = R.bid AND B.color = 'red'
UNION ALL
SELECT R.sid
FROM Boats B, Reserves R
WHERE B.bid = R.bid AND B.color = 'blue'
```

sid of Sailors that reserved red or blue boat

```
SELECT DISTINCT R.sid
FROM Boats B, Reserves R
WHERE B.bid = R.bid AND
      (B.color = 'red' OR B.color = 'blue')
```

OR

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE B.bid = R.bid AND B.color = 'red'
UNION
SELECT R.sid
FROM Boats B, Reserves R
WHERE B.bid = R.bid AND B.color = 'blue'
```

sid of Sailors that reserved red and blue boat

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE B.bid = R.bid AND
      (B.color = 'red' AND B.color = 'blue')
```

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE B.bid = R.bid AND B.color = 'red'
INTERSECT ALL
SELECT R.sid
FROM Boats B, Reserves R
WHERE B.bid = R.bid AND B.color = 'blue'
```

sid of Sailors that reserved red and blue boat

Can use self-join instead

```
SELECT R.sid
FROM Boats B1, Reserves R1
WHERE B1.bid = R1.bid AND
      B1.color = 'red'
```

sid of Sailors that reserved red and blue boat

Can use self-join instead

```
SELECT R.sid
FROM Boats B1, Reserves R1, Boats B2, Reserves R2
WHERE B1.bid = R1.bid AND
      B1.color = 'red'
```

sid of Sailors that reserved red and blue boat

Can use self-join instead

```
SELECT R.sid
FROM Boats B1, Reserves R1, Boats B2, Reserves R2
WHERE B1.bid = R1.bid AND
      B2.bid = R2.bid AND
      B1.color = 'red' AND B2.color = 'blue'
```

sid of Sailors that reserved red and blue boat

Can use self-join instead

```
SELECT R.sid
FROM Boats B1, Reserves R1, Boats B2, Reserves R2
WHERE R1.sid = R2.sid AND
      B1.bid = R1.bid AND
      B2.bid = R2.bid AND
      B1.color = 'red' AND B2.color = 'blue'
```

sids of sailors that haven't reserved a boat

```
SELECT S.sid
FROM Sailors S

EXCEPT

SELECT S.sid
FROM Sailors S, Reserves R
WHERE S.sid = R.sid
```

Can we write EXCEPT using more basic functionality?

## SET Comparison Operators

UNION, INTERSECT, EXCEPT

EXISTS, NOT EXISTS  
IN, NOT IN  
UNIQUE, NOT UNIQUE

*op* ANY, *op* ALL  
 $op \in \{<, >, =, \leq, \geq, \neq, \dots\}$

Many of these rely on Nested Query Support

## Nested Queries

```
SELECT S.sid
FROM Sailors S
WHERE S.sid IN (SELECT R.sid
                FROM Reserves R
                WHERE R.bid = 101)
```

Many clauses can contain SQL queries  
WHERE, FROM, HAVING, SELECT

Conceptual model:  
for each Sailors tuple  
run the subquery and evaluate qualification

## Nested Correlated Queries

```
SELECT S.sid
FROM Sailors S
WHERE EXISTS (SELECT *
              FROM Reserves R
              WHERE R.bid = 101 AND
                  S.sid = R.sid)
```

Outer table referenced in nested query

Conceptual model:  
for each Sailors tuple  
run the subquery and evaluate qualification

## Nested Correlated Queries

```
SELECT S.sid
FROM Sailors S
WHERE UNIQUE (SELECT *
              FROM Reserves R
              WHERE R.bid = 101 AND
                  S.sid = R.sid)
```

UNIQUE checks that there are no duplicates

What does this do?

## Nested Correlated Queries

```
SELECT S.sid
FROM Sailors S
WHERE UNIQUE (SELECT R.sid
              FROM Reserves R
              WHERE R.bid = 101 AND
                  S.sid = R.sid)
```

UNIQUE checks that there are no duplicates

What does this do?

Sailors whose rating is greater than  
any sailor named "Bobby"

```
SELECT S1.name
FROM Sailors S1
WHERE S1.rating > ANY (SELECT S2.rating
                        FROM Sailors S2
                        WHERE S2.name = 'Bobby')
```

What about this?

```
SELECT S1.name
FROM Sailors S1
WHERE S1.rating > ALL (SELECT S2.rating
                       FROM Sailors S2
                       WHERE S2.name = 'Bobby')
```

Rewrite INTERSECT using IN

```
SELECT S.sid
FROM Sailors S
WHERE S.rating > 2
INTERSECT
SELECT R.sid
FROM Reserves R
)
SELECT S.sid
FROM Sailors S
WHERE S.rating > 2 AND
      S.sid IN (
        SELECT R.sid
        FROM Reserves R
      )
```

Similar trick for EXCEPT → NOT IN

What if want *names* instead of sids?

Sailors that reserved all boats (Division)

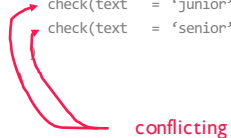
Hint: double negation  
reserved all boats == no boat w/out reservation

```
SELECT S.name
FROM Sailors S
WHERE NOT EXISTS (
  (SELECT B.bid FROM Boats B)
  EXCEPT
  (SELECT R.bid
   FROM Reserves R
   WHERE R.sid = S.sid)
)
```

HWI bugs

Conflicting CHECK constraints

```
Prof(
  type text,
  check(text in ('junior', 'senior')),
  check(text = 'junior' and hired is not null),
  check(text = 'senior' and tenure_year is not null)
```



conflicting

HWI bugs

At most once *per semester* translated as at most once

```
CREATE TABLE Offers (
  deptid text,
  courseid text,
  semester text,
  year int,
  . . .
  PRIMARY KEY(deptid, courseid)
);
```

Wrong

## HWI bugs

At most once *per semester* translated as at most once

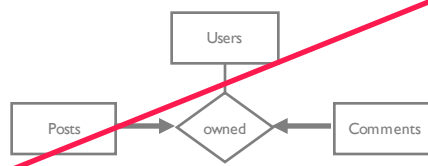
```
CREATE TABLE Offers (
  deptid text,
  courseid text,
  semester text,
  year int,
  . . .
  PRIMARY KEY(deptid, courseid, semester, year)
);
```

## HWI bugs

Reddit:

Comments owned by one user

Posts owned by one user




## HWI bugs

Reddit:

Comments owned by one user

Posts owned by one user



## Sailors that reserved all boats (Division)

Hint: double negation

reserved all boats == no boat w/out reservation

```
SELECT S.name
FROM Sailors S
WHERE NOT EXISTS (
```

Sailors S such that

There's no boat without

A reservation by S

## Sailors that reserved all boats (Division)

Hint: double negation

reserved all boats == no boat w/out reservation

```
SELECT S.name
FROM Sailors S
WHERE NOT EXISTS (SELECT B.bid
                  FROM Boats B
                  WHERE NOT EXISTS (
```

Sailors S such that

There's no boat without

A reservation by S

## Sailors that reserved all boats (Division)

Hint: double negation

reserved all boats == no boat w/out reservation

```
SELECT S.name
FROM Sailors S
WHERE NOT EXISTS (SELECT B.bid
                  FROM Boats B
                  WHERE NOT EXISTS (SELECT R.bid
                                    FROM Reserves R
                                    WHERE R.sid = S.sid))
```

Sailors S such that

There's no boat without

A reservation by S



## NULL

Field values sometimes unknown or inapplicable  
SQL provides a special value *null* for such situations.

The presence of null complicates many issues e.g.,

Is age = null true or false?

Is null = null true or false?

Is null = 8 OR 1 = 1 true or false?

Special syntax "IS NULL" and "IS NOT NULL"  
3 Valued Logic (true, false, unknown)

How does WHERE remove rows?

if qualification doesn't evaluate to true

New operators (in particular, outer joins) possible/needed.

## NULL

(null > 0) = null

(null + 1) = null

(null = 0) = null

(null AND true) = null

null is null = true

### Some truth tables

AND	T	F	NULL
T	T	F	NULL
F	F	F	F
NULL	NULL	F	NULL

OR	T	F	NULL
T	T	T	T
F	T	F	NULL
NULL	T	NULL	NULL

## JOINS

```
SELECT [DISTINCT] target_List
FROM table_name
    [INNER | {LEFT | RIGHT | FULL } {OUTER}] JOIN table_name
    ON qualification_List
WHERE ...
```

INNER is default

Difference in how to deal with NULL values

PostgreSQL documentation:

<http://www.postgresql.org/docs/9.4/static/tutorial-join.html>

## Inner/Natural Join

```
SELECT s.sid, s.name, r.bid
FROM Sailors S, Reserves r
WHERE s.sid = r.sid
```

```
SELECT s.sid, s.name, r.bid
FROM Sailors s INNER JOIN Reserves r
ON s.sid = r.sid
```

All  
Equivalent!

```
SELECT s.sid, s.name, r.bid
FROM Sailors s NATURAL JOIN Reserves r
```

**Natural Join** means equi-join for each pair of  
attrs with same name

### Sailor names and their reserved boat ids

```
SELECT s.sid, s.name, r.bid
FROM Sailors s INNER JOIN Reserves r
ON s.sid = r.sid
```

Sailors

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Reserves

sid	bid	day
1	102	9/12
2	102	9/13

Result

sid	name	bid
1	Eugene	102
2	Luis	102

### Sailor names and their reserved boat ids

```
SELECT s.sid, s.name, r.bid
FROM Sailors s INNER JOIN Reserves r
ON s.sid = r.sid
```

Sailors

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Reserves

sid	bid	day
1	102	9/12
2	102	9/13

Result

sid	name	bid
1	Eugene	102
2	Luis	102

Prefer INNER JOIN over NATURAL JOIN. Why?

### Sailor names and their reserved boat ids

```
SELECT s.sid, s.name, r.bid
FROM Sailors s INNER JOIN Reserves r
ON s.sid = r.sid
```

Sailors

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Reserves

sid	bid	day
1	102	9/12
2	102	9/13

Result

sid	name	bid
1	Eugene	102
2	Luis	102

Notice: No result for Ken!

### Left Outer Join (or No Results for Ken)

Returns all matched rows *and all unmatched rows from table on left of join clause*

(at least one row for each row in left table)

```
SELECT s.sid, s.name, r.bid
FROM Sailors s LEFT OUTER JOIN Reserves r
ON s.sid = r.sid
```

All sailors & bid for boat in their reservations  
Bid set to NULL if no reservation

### Left Outer Join

```
SELECT s.sid, s.name, r.bid
FROM Sailors s LEFT OUTER JOIN Reserves r
ON s.sid = r.sid
```

Sailors

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Reserves

sid	bid	day
1	102	9/12
2	102	9/13

Result

sid	name	bid
1	Eugene	102
2	Luis	102
3	Ken	NULL

### Can Left Outer Join be expressed with Cross-Product?

Sailors

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Reserves

sid	bid	day
-----	-----	-----

Sailors x Reserves

Sailors s **LEFT OUTER JOIN** Reserves r  
ON s.sid = r.sid

Result

sid	name	bid
-----	------	-----

Result

sid	name	bid
1	Eugene	NULL
2	Luis	NULL
3	Ken	NULL

### Can Left Outer Join be expressed with Cross-Product?

Sailors

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Reserves

sid	bid	day
-----	-----	-----

Sailors  $\bowtie$  Reserves

U

(Sailors - (Sailors  $\bowtie$  Reserves)) x {(null, ... )}

How to compute this with a query?

### Right Outer Join

Same as LEFT OUTER JOIN, but guarantees result for rows in table on *right side of JOIN*

```
SELECT s.sid, s.name, r.bid
FROM Reserves r RIGHT OUTER JOIN Sailors s
ON s.sid = r.sid
```

## FULL OUTER JOIN

Returns all matched or unmatched rows from both sides of JOIN

```
SELECT s.sid, s.name, r.bid
FROM Sailors s FULL OUTER JOIN Reserves r
ON s.sid = r.sid
```

## FULL OUTER JOIN

```
SELECT s.sid, s.name, r.bid
FROM Sailors s FULL OUTER JOIN Reserves r
ON s.sid = r.sid
```

Sailors

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Reserves

sid	bid	day
1	102	9/12
2	102	9/13
4	109	9/20

Result

sid	name	bid
1	Eugene	102
2	Luis	102
3	Ken	NULL
NULL	NULL	109

Why is sid NULL?

## Serious people can count: Aggregation

```
SELECT COUNT(*)
FROM Sailors S
COUNT([DISTINCT] A)

SELECT AVG(S.age)
FROM Sailors S
WHERE S.rating = 10
SUM([DISTINCT] A)

SELECT COUNT(DISTINCT S.name)
FROM Sailors S
WHERE S.name LIKE 'D%'
AVG([DISTINCT] A)
MAX/MIN(A)
STDDEV(A)
CORR(A,B)
```

PostgreSQL documentation  
<http://www.postgresql.org/docs/9.4/static/functions-aggregate.html>

## Name and age of oldest sailor(s)

```
SELECT S.name, MAX(S.age)
FROM Sailors S
```

```
SELECT S.name, S.age
FROM Sailors S
WHERE S.age >= ALL (SELECT S2.age
FROM Sailors S2)
```

```
SELECT S.name, S.age
FROM Sailors S
WHERE S.age = (SELECT MAX(S2.age)
FROM Sailors S2)
```

```
SELECT S.name, S.age
FROM Sailors S
ORDER BY S.age DESC
LIMIT 1
```

← When does this not work?

## GROUP BY

```
SELECT min(s.age)
FROM Sailors s
```

Minimum age among all sailors

What if want min age *per rating level*?

We don't even know how many rating levels exist!

If we did, could write (awkward):

```
for rating in [0..10]
SELECT min(s.age)
FROM Sailors s
WHERE s.rating = <rating>
```

## GROUP BY

```
SELECT count(*)
FROM Reserves R
```

Total number of reservations

What if want reservations per boat?

May not even know all our boats (depends on data)!

If we did, could write (awkward):

```
for boat in [0..10]
SELECT count(*)
FROM Reserves R
WHERE R.bid = <boat>
```

## GROUP BY

```

SELECT    [DISTINCT] target-list
FROM      relation-list
WHERE     qualification
GROUP BY  grouping-list
HAVING    group-qualification

```

*grouping-list* is a list of expressions that defines groups  
set of tuples w/ same value for all attributes in *grouping-list*

*target-list* contains


$attribute-names \subseteq grouping-list$   
aggregation expressions

## Conceptual Query Evaluation

```

SELECT    [DISTINCT] target-list
FROM      relation-list
WHERE     qualification
GROUP BY  grouping-list
HAVING    group-qualification

```

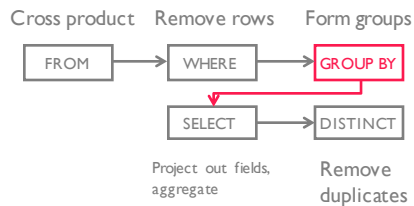


## Conceptual Query Evaluation

```

SELECT    [DISTINCT] target-list
FROM      relation-list
WHERE     qualification
GROUP BY  grouping-list
HAVING    group-qualification

```

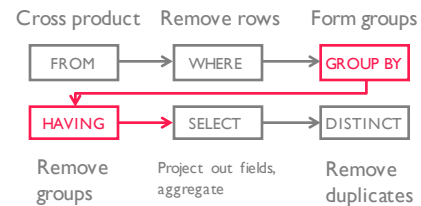


## Conceptual Query Evaluation

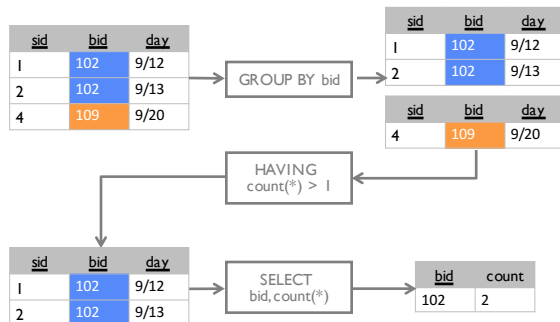
```

SELECT    [DISTINCT] target-list
FROM      relation-list
WHERE     qualification
GROUP BY  grouping-list
HAVING    group-qualification

```



## Conceptual Evaluation



## GROUP BY

```

SELECT    bid, count(*)
FROM      Reserves R
GROUP BY  bid

```

Minimum age for each rating

```

SELECT    bid, count(*)
FROM      Reserves R
GROUP BY  bid
HAVING    count(*) > 1

```

Minimum age for each boat with  
more than 1 reservation

## HAVING

*group-qualification* used to remove groups  
similar to WHERE clause

Expressions must have *one value per group*. Either  
An aggregation function or  
In *grouping-list*

```
SELECT bid, count(*)
FROM reserves R
GROUP BY bid
HAVING color = 'red'
```

## AVG age of sailors reserving red boats, by rating

```
SELECT
FROM Sailors S, Boats B, Reserves R
WHERE S.sid = R.sid AND
      R.bid = B.bid AND
      B.color = 'red'
```

## AVG age of sailors reserving red boats, by rating

```
SELECT S.rating, avg(S.age) AS age
FROM Sailors S, Boats B, Reserves R
WHERE S.sid = R.sid AND
      R.bid = B.bid AND
      B.color = 'red'
GROUP BY S.rating
```

What if move B.color='red' to HAVING clause?

## Ratings where the avg age is min over all ratings



```
SELECT S.rating
FROM Sailors S
WHERE S.age = (
  SELECT MIN(AVG(S2.age))
  FROM Sailors S2
)
```



```
SELECT S.rating
FROM (SELECT S.rating, AVG(S.age) as avgage
      FROM Sailors S
      GROUP BY S.rating) AS tmp
WHERE tmp.avgage = (
  SELECT MIN(tmp2.avgage) FROM (
    SELECT S.rating, AVG(S.age) as avgage
    FROM Sailors S
    GROUP BY S.rating
  ) AS tmp2
)
```

## Ratings where the avg age is min over all ratings



```
SELECT S.rating
FROM Sailors S
WHERE S.age = (
  SELECT MIN(AVG(S2.age))
  FROM Sailors S2
)
```



```
SELECT S.rating
FROM (SELECT S.rating, AVG(S.age) as avgage
      FROM Sailors S
      GROUP BY S.rating) AS tmp
WHERE tmp.avgage <= ALL (
  SELECT tmp2.avgage FROM (
    SELECT S.rating, AVG(S.age) as avgage
    FROM Sailors S
    GROUP BY S.rating
  ) AS tmp2
)
```

## Setting up Proj I Part 2

Users assigned to schemas (namespaces).  
Noted user didn't have an assigned schema  
User created their tables under Public schema.  
Uh oh! Did I miss anyone else

Students without a schema

username
vx3948
et1827
etu4938

```
SELECT username
FROM pg_user;
```

schemaname
et2039
sa2037
kt6765

```
SELECT schemaname
FROM pg_tables;
```

## Setting up Proj I Part 2

```
FROM (SELECT username FROM pg_user) AS U
      (SELECT schemaname FROM pg_tables) AS S
```

username
vx3948
et1827
etu4938

```
SELECT username
FROM pg_user;
```

schemaname
et2039
sa2037
kt6765

```
SELECT schemaname
FROM pg_tables;
```

## Setting up Proj I Part 2

```
FROM (SELECT username FROM pg_user) AS U
      RIGHT OUTER JOIN
      (SELECT schemaname FROM pg_tables) AS S
      ON U.username = S.schemaname
```

username
vx3948
et1827
etu4938

```
SELECT username
FROM pg_user;
```

schemaname
et2039
sa2037
kt6765

```
SELECT schemaname
FROM pg_tables;
```

## Setting up Proj I Part 2

```
SELECT U.username, S.schemaname
FROM (SELECT username FROM pg_user) AS U
      RIGHT OUTER JOIN
      (SELECT schemaname FROM pg_tables) AS S
      ON U.username = S.schemaname
WHERE S.schemaname is null;
```

username
vx3948
et1827
etu4938

```
SELECT username
FROM pg_user;
```

schemaname
et2039
sa2037
kt6765

```
SELECT schemaname
FROM pg_tables;
```


## Setting up Proj I Part 2

```
SELECT count(*)
FROM (
  SELECT U.username, S.schemaname
  FROM (SELECT username FROM pg_user) AS U
        RIGHT OUTER JOIN
        (SELECT schemaname FROM pg_tables) AS S
        ON U.username = S.schemaname
  WHERE S.schemaname is null;
) AS nestedquery
```

124  
but 190 users!

## ORDER BY, LIMIT

```
SELECT [DISTINCT] target-List
FROM relation-List
WHERE qualification
GROUP BY grouping-List
HAVING group-qualification
ORDER BY order-List
LIMIT Limit-expr [OFFSET offset-expr]
```



## ORDER BY

```
SELECT S.name
FROM Sailors S
ORDER BY (S.rating/2)::int ASC,
         S.age DESC
```

List of *order-list* expressions dictates ordering precedence  
Sorted in ascending by age/rating ratio  
If ties, sorted high to low rating

## ORDER BY

```
SELECT  S.name, (S.rating/2)::int, S.age
FROM    Sailors S
ORDER BY (S.rating/2)::int ASC,
         S.age DESC
```

Sailors

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Result

name	int4	age
Luis	1	39
Ken	4	27
Eugene	4	22

## ORDER BY

```
SELECT  S.name, (S.rating/2)::int, S.age
FROM    Sailors S
ORDER BY (S.rating/2)::int ASC,
         S.age ASC
```

Sailors

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Result

name	int4	age
Luis	1	39
Eugene	4	22
Ken	4	27

## LIMIT

```
SELECT  S.name, (S.rating/2)::int, S.age
FROM    Sailors S
ORDER BY (S.rating/2)::int ASC,
         S.age DESC
LIMIT 2
```

Only the first 2 results

Sailors

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Result

name	int4	age
Luis	1	39
Ken	4	27

## LIMIT

```
SELECT  S.name, (S.rating/2)::int, S.age
FROM    Sailors S
ORDER BY (S.rating/2)::int ASC,
         S.age DESC
LIMIT 2 OFFSET 1
```

Only the first 2 results

Sailors

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Result

name	int4	age
Ken	4	27
Eugene	4	22

## LIMIT

```
SELECT  S.name, (S.rating/2)::int, S.age
FROM    Sailors S
ORDER BY (S.rating/2)::int ASC,
         S.age DESC
LIMIT (SELECT count(S2.*) / 2
       FROM Sailors AS S2)
```

Can have expressions instead of constants

Result

name	int4	age
Luis	1	39

## Integrity Constraints

Conditions that every legal instance must satisfy  
 Inserts/Deletes/Updates that violate ICs rejected  
 Helps ensure app semantics or prevent inconsistencies

We've discussed  
 domain/type constraints, primary/foreign key  
 general constraints

## Beyond Keys: Table Constraints

Runs when table is not empty

```
CREATE TABLE Sailors(
  sid int,
  ...
  PRIMARY KEY (sid),
  CHECK (rating >= 1 AND rating <= 10)
```

```
CREATE TABLE Reserves(
  sid int,
  bid int,
  day date,
  PRIMARY KEY (bid, day),
  CONSTRAINT no_red_reservations
  CHECK ('red' <> (SELECT B.color
                    FROM Boats B
                    WHERE B.bid = bid))
```

Nested subqueries  
Named constraints

## Multi-Relation Constraints

# of sailors + # of boats should be less than 100

```
CREATE TABLE Sailors (
  sid int,
  bid int,
  day date,
  PRIMARY KEY (bid, day),
  CHECK (
    (SELECT COUNT(S.sid) FROM Sailors S)
    +
    (SELECT COUNT(B.bid) FROM Boats B)
    < 100
  )
```

What if Sailors is empty?

## ASSERTIONS: Multi-Relation Constraints

```
CREATE ASSERTION small_club
CHECK (
  (SELECT COUNT(*) FROM Sailors S)
  +
  (SELECT COUNT(*) FROM Boats B)
  < 100
)
```

**ASSERTIONS** are not associated with any table

## WHAT!

So many things we can't express or don't work!

Assertions

Nested queries in CHECK constraints



## Advanced Stuff

User defined functions

Triggers

WITH

Views

## User Defined Functions (UDFs)

Custom functions that can be called in database

Many languages: SQL, python, C, perl, etc

```
CREATE FUNCTION function_name(p1 type, p2 type, ...)
RETURNS type
```



## User Defined Functions (UDFs)

Custom functions that can be called in database

Many languages: SQL, python, C, perl, etc

```
CREATE FUNCTION function_name(p1 type, p2 type, ...)
RETURNS type
AS $$
```

```
-- Logic
```

```
$$ LANGUAGE language_name;
```

## User Defined Functions (UDFs)

Custom functions that can be called in database

Many languages: SQL, python, C, perl, etc

```
CREATE FUNCTION function_name(p1 type, p2 type, ...)
RETURNS type
AS $$
```

```
-- Logic
```

```
$$ LANGUAGE language_name;
```

## A simple UDF (lang = SQL)

```
CREATE FUNCTION mult1(v int) RETURNS int
AS $$
SELECT v * 100;
$$ LANGUAGE SQL;
```

← Last statement  
is returned

```
CREATE FUNCTION function_name(p1 type, p2 type, ...)
RETURNS type
AS $$
```

```
-- Logic
```

```
$$ LANGUAGE language_name;
```

## A simple UDF (lang = SQL)

```
CREATE FUNCTION mult1(v int) RETURNS int
AS $$
SELECT v * 100;
$$ LANGUAGE SQL;
```

```
SELECT mult1(S.age)
FROM sailors AS S
```

Sailors

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Result

int4
220
390
270

<http://www.postgresql.org/docs/9.1/func-sql.html>

## A simple UDF (lang = SQL)

```
CREATE FUNCTION mult1(v int) RETURNS int
AS $$
SELECT $1 * 100;
$$ LANGUAGE SQL;
```

```
SELECT mult1(S.age)
FROM sailors AS S
```

Sailors

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Result

int4
220
390
270

<http://www.postgresql.org/docs/9.1/func-sql.html>

## Process a Record (lang = SQL)

```
CREATE FUNCTION mult2(x sailors) RETURNS int
AS $$
SELECT (x.sid + x.age) / x.rating;
$$ LANGUAGE SQL;
```

```
SELECT mult2(S.*)
FROM sailors AS S
```

Sailors

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Result

int4
3.285
20.5
3.75

<http://www.postgresql.org/docs/9.1/func-sql.html>

## Process a Record (lang = SQL)

```
CREATE FUNCTION mult2(sailors) RETURNS int
AS $$
SELECT ($1.sid + $1.age) / $1.rating;
$$ LANGUAGE SQL;
```

```
SELECT mult2(S.*)
FROM sailors AS S
```

Sailors

sid	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Result

int4
3.285
20.5
3.75

<http://www.postgresql.org/docs/9.4/static/func-sql.html>

## Procedural Language/SQL (lang = plsql)

```
CREATE FUNCTION proc(v int) RETURNS int
AS $$
DECLARE
    -- define variables
BEGIN
    -- PL/SQL code
END;
$$ LANGUAGE plpgsql;
```

Boilerplate

<http://www.postgresql.org/docs/9.4/static/plpgsql.html>

## Procedural Language/SQL (lang = plsql)

```
CREATE FUNCTION proc(v int) RETURNS int
AS $$
DECLARE
    -- define variables.  VAR TYPE [= value]
    qty int = 10;
BEGIN
    qty = qty * v;
    INSERT INTO blah VALUES(qty);
    RETURN qty + 2;
END;
$$ LANGUAGE plpgsql;
```

<http://www.postgresql.org/docs/9.4/static/plpgsql.html>

## Procedural Code (lang = plpython2u)

```
CREATE FUNCTION proc(v int) RETURNS int
AS $$
import random
return random.randint(0, 100) * v
$$ LANGUAGE plpython2u;
```

Very powerful – can do anything so must be careful

run in a python interpreter with no security protection

plpy module provides database access

```
plpy.execute("select 1")
```

<http://www.postgresql.org/docs/9.4/static/plpython.html>

## Procedural Code (lang = plpython2u)

```
CREATE FUNCTION proc(word text) RETURNS text
AS $$
import requests
resp = requests.get('http://google.com/search?q=%s' % v)
return resp.content
$$ LANGUAGE plpython2u;
```

Very powerful – can do anything so must be careful

run in a python interpreter with no security protection

plpy module provides database access

```
plpy.execute("select 1")
```

<http://www.postgresql.org/docs/9.4/static/plpython.html>

## Triggers (logical)

def: procedure that runs automatically if specified changes in DBMS happen

```
CREATE TRIGGER name
```

Event activates the trigger

Condition tests if triggers should run

Action what to do

## Triggers (logical)

def: procedure that runs automatically if specified changes in DBMS happen

```
CREATE TRIGGER name
[BEFORE | AFTER | INSTEAD OF] event_list
ON table
```

**Event** activates the trigger

**Condition** tests if triggers should run

**Action** what to do

## Triggers (logical)

def: procedure that runs automatically if specified changes in DBMS happen

```
CREATE TRIGGER name
[BEFORE | AFTER | INSTEAD OF] event_list
ON table

WHEN trigger_qualifications
```

**Event** activates the trigger

**Condition** tests if triggers should run

**Action** what to do

## Triggers (logical)

def: procedure that runs automatically if specified changes in DBMS happen

```
CREATE TRIGGER name
[BEFORE | AFTER | INSTEAD OF] event_list
ON table
[FOR EACH ROW]
WHEN trigger_qualifications
procedure
```

**Event** activates the trigger

**Condition** tests if triggers should run

**Action** what to do

## Copy new young sailors into special table (logical)

```
CREATE TRIGGER youngSailorUpdate
AFTER INSERT ON SAILORS
REFERENCING NEW TABLE NewInserts
FOR EACH STATEMENT
INSERT
    INTO YoungSailors(sid, name, age, rating)
    SELECT sid, name, age, rating
    FROM NewInserts N
    WHERE N.age <= 18
```

**Event** activates the trigger

**Condition** tests if triggers should run

**Action** what to do

## Copy new young sailors into special table (logical)

```
CREATE TRIGGER youngSailorUpdate
AFTER INSERT ON SAILORS
FOR EACH ROW
WHEN NEW.age <= 18
INSERT
    INTO YoungSailors (sid, name, age, rating)
    VALUES (NEW.sid, NEW.name, NEW.age, NEW.rating)
```

**Event** activates the trigger

**Condition** tests if triggers should run

**Action** what to do

## Triggers (logical)

Can be complicated to reason about

Triggers may (e.g., insert) cause other triggers to run  
If > 1 trigger match an action, which is run first?

╰(ˊˋ)╯

```
CREATE TRIGGER recursiveTrigger
AFTER INSERT ON SAILORS
FOR EACH ROW
INSERT INTO Sailors(sid, name, age, rating)
SELECT sid, name, age, rating
FROM Sailors S
```

## Triggers (postgres)

```
CREATE TRIGGER name
[BEFORE | AFTER | INSTEAD OF] event_list
ON table
FOR EACH (ROW | STATEMENT)
WHEN trigger_qualifications
EXECUTE PROCEDURE user_defined_function();
```

PostgreSQL only runs *trigger* UDFs

<http://www.postgresql.org/docs/9.1/static/sql-createtrigger.html>  
<http://www.postgresql.org/docs/9.1/static/plpgsql-trigger.html>

## Trigger Example

```
CREATE FUNCTION copyrecord() RETURNS trigger
AS $$
BEGIN
    INSERT INTO blah VALUES(NEW.a);
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;
```

Signature: no args, return type is trigger

Returns NULL or same record structure as modified row

Special variables: OLD, NEW

```
CREATE TRIGGER t_copyinserts BEFORE INSERT ON a
FOR EACH ROW
EXECUTE PROCEDURE copyrecord();
```

<http://www.postgresql.org/docs/9.1/static/sql-createtrigger.html>  
<http://www.postgresql.org/docs/9.1/static/plpgsql-trigger.html>

## Total boats and sailors < 100

```
CREATE FUNCTION checktotal() RETURNS trigger
AS $$
BEGIN
    IF ((SELECT COUNT(*) FROM sailors) +
        (SELECT COUNT(*) FROM boats) < 100) THEN
        RETURN NEW
    ELSE
        RETURN null;
    END IF;
END;
$$ LANGUAGE plpgsql;
```

```
CREATE TRIGGER t_checktotal BEFORE INSERT ON sailors
FOR EACH ROW
EXECUTE PROCEDURE checktotal();
```

<http://www.postgresql.org/docs/9.1/static/sql-createtrigger.html>  
<http://www.postgresql.org/docs/9.1/static/plpgsql-trigger.html>

## You can get into trouble...

```
CREATE FUNCTION addme_bad() RETURNS trigger
AS $$
BEGIN
    INSERT INTO a VALUES (NEW.*);
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;
```

```
CREATE TRIGGER t_addme_bad BEFORE INSERT ON a
FOR EACH ROW
EXECUTE PROCEDURE addme_bad();
```

<http://www.postgresql.org/docs/9.1/static/sql-createtrigger.html>  
<http://www.postgresql.org/docs/9.1/static/plpgsql-trigger.html>

## You can get into trouble...

```
CREATE FUNCTION addme_ok() RETURNS trigger
AS $$
BEGIN
    IF (SELECT COUNT(*) FROM a) < 100 THEN
        INSERT INTO a VALUES (NEW.a + 1);
    END IF;
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;
```

```
CREATE TRIGGER t_addme_ok BEFORE INSERT ON a
FOR EACH ROW
EXECUTE PROCEDURE addme_ok();
```

<http://www.postgresql.org/docs/9.1/static/sql-createtrigger.html>  
<http://www.postgresql.org/docs/9.1/static/plpgsql-trigger.html>

## You can get into trouble...

```
CREATE FUNCTION addme_works() RETURNS trigger
AS $$
BEGIN
    IF (SELECT COUNT(*) FROM a) < 100 THEN
        INSERT INTO a VALUES (NEW.a + 1);
    END IF;
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;
```

```
CREATE TRIGGER t_addme_works AFTER INSERT ON a
FOR EACH ROW
EXECUTE PROCEDURE addme_works();
```

<http://www.postgresql.org/docs/9.1/static/sql-createtrigger.html>  
<http://www.postgresql.org/docs/9.1/static/plpgsql-trigger.html>

## WITH

```
WITH RedBoats(bid, count) AS
(SELECT B.bid, count(*)
 FROM Boats B, Reserves R
 WHERE R.bid = B.bid AND B.color = 'red'
 GROUP BY B.bid)
SELECT name, count
FROM Boats AS B, RedBoats AS RB
WHERE B.bid = RB.bid AND count < 2
```

Names of unpopular boats

## Views

```
CREATE VIEW view_name
AS select_statement
```

“tables” defined as query results rather than inserted base data

Makes development simpler

Used for security

Not materialized

References to *view\_name* replaced with *select\_statement*

Similar to WITH, lasts longer than one query

## Views

```
CREATE VIEW boat_counts
AS SELECT bid, count(*)
FROM Reserves R
GROUP BY bid
HAVING count(*) > 10
```

Used like a normal table

```
SELECT bname
FROM boat_counts bc, Boats B
WHERE bc.bid = B.bid
```

Names of popular boats

```
SELECT bname
FROM
  (SELECT bid, count(*)
   FROM Reserves R
   GROUP BY bid
   HAVING count(*) > 10) bc,
  Boats B
WHERE bc.bid = B.bid
```

Rewritten expanded query

## CREATE TABLE

```
CREATE TABLE <table_name> AS
<SELECT STATEMENT>
```

Guess the schema:

```
CREATE TABLE used_boats1 AS
SELECT r.bid
FROM Sailors s,
     Reservations r
WHERE s.sid = r.sid

CREATE TABLE used_boats2 AS
SELECT r.bid as foo
FROM Sailors s,
     Reservations r
WHERE s.sid = r.sid

used_boats1(bid int)
used_boats2(foo int)
```

How is this different than views?

What if we insert a new record into Reservations?

## Summary

SQL is pretty complex

Superset of Relational Algebra SQL99 turing complete

Human readable

More than one way to skin a horse

Many alternatives to write a query

Optimizer (theoretically) finds most efficient plan

