

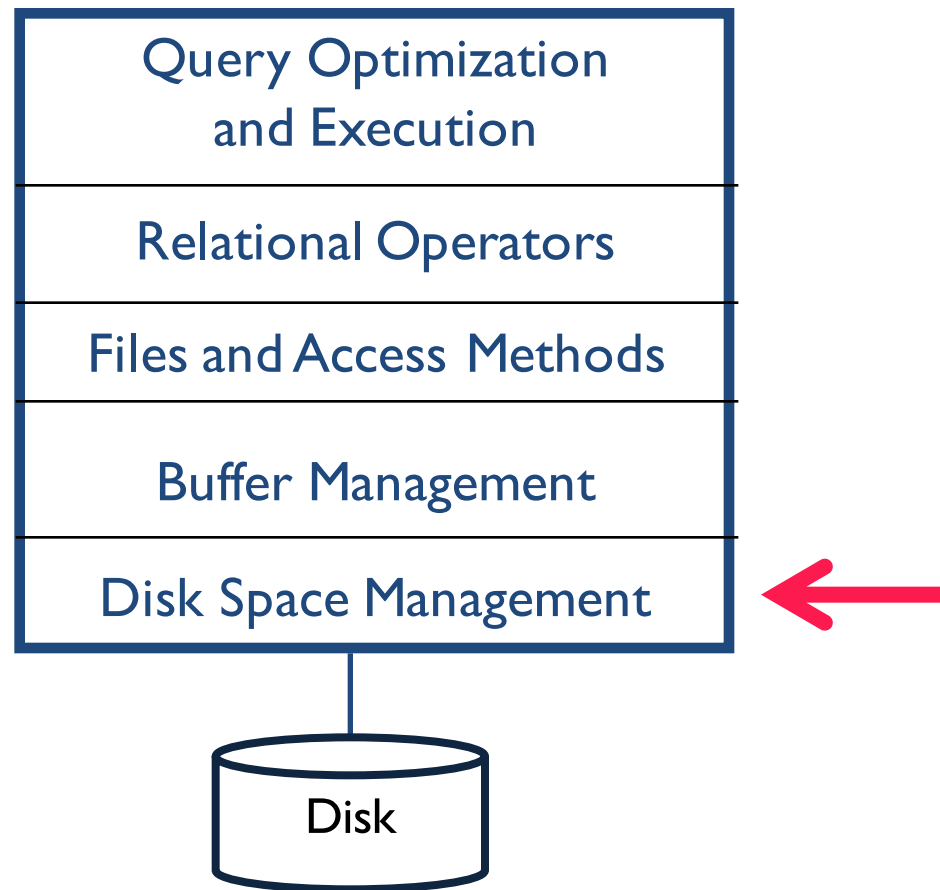
**L8**

# **Disk, Storage, and Indexing**

Eugene Wu

Fall 2015

# Work from the bottom up



# \$ Matters

Why not store all in RAM?

Costs too much

High-end Databases today ~Petabyte (1000TB) range.  
~60% cost of a production system is in the disks.

Main memory not persistent

Obviously important if DB stops/crashes

Some systems are *main-memory* DBMSes, topic for advanced DB course


# \$ Matters

Newegg enterprise \$1000

RAM: 64

SSD: 400

Disk: 4000



RAM for active data

Registers

.3 ns

Disk for main database  
secondary storage

on chip cache

.5ns

off chip cache

7 ns

Tapes for archive

RAM

100 ns

SSD

.15 ms

SSD is still in flux,  
some in use

Disk


5 ms

Interesting numbers

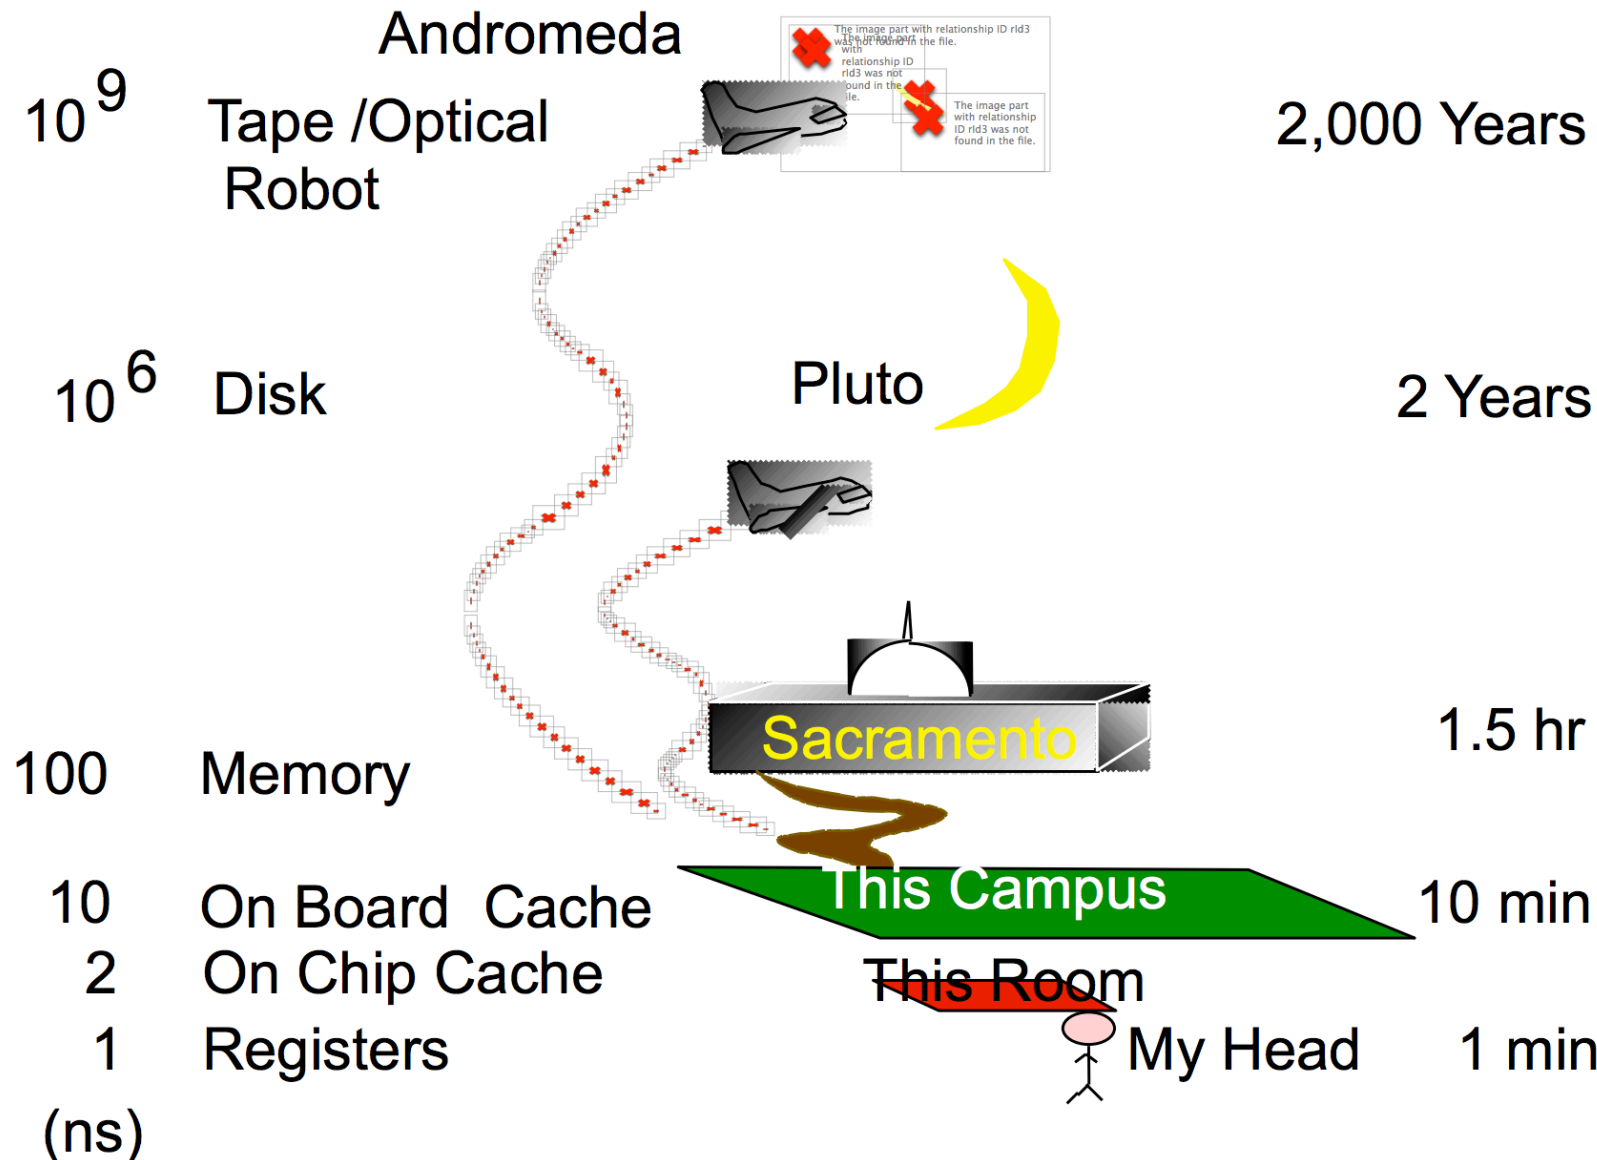
compress 1k bytes: 3000 ns

roundtrip in data center: .5 ms

# Jim Gray's Storage Latency Analogy: How Far Away is the Data?




**Jim Gray**  
Turing Award  
B.S. Cal 1966  
Ph.D. Cal 1969!



Spin speed: ~7200 RPM

Arm moved in/out to  
position head over a  
track



Time to access (read or write) a disk block

seek time                      2-4 msec avg

rotational delay              2-4 msec

transfer time                 0.3 msec/64kb page

Throughput

read                            ~150 MB/sec

write                           ~50 MB/sec

Key: reduce seek and rotational delays

HW & SW approaches



Next block concept (in order of speed)

- blocks on same track

- blocks on same cylinder

- blocks on adjacent cylinder

Sequentially arrange files

- minimize seek and rotation latency

When sequentially scanning: Pre-fetch

- > 1 page/block at once

# SSD maybe

Fast changing, not yet stabilized

Read small & fast

single read:	0.03ms
4kb random reads:	500MB/sec
seq reads:	525MB/sec

Write is slower for random


single write:	0.03ms
4kb random writes:	120MB/sec
seq writes:	120MB/sec

Write endurance limited

2-3k cycle lifetimes

6-10 months

# # 4 byte values read per second



5 orders of magnitude

# Pragmatics of Databases

Most databases are pretty small

All global daily weather since 1929: 20GB

2000 US Census: 200GB

2009 english wikipedia: 14GB

Data sizes grow faster than moore's law

# Disk Space Management

VLDBs                      SSDs: reduce variance  
Small DBs                interesting data is small


Huge data exists

Many interesting data is small

People will still worry about magnetic disk.

May not care about it

# Work from the bottom up



# Disk Space Management

Lowest layer of DBMS, manages space on disk

Low level IO interface:

- allocate/deallocate a page

- read/write page

Sequential performance desirable

- try to ensure sequential pages are sequential on disk

- hidden from rest of DBMS

- but algorithms may assume sequential performance

# Files

Pages are IO interface

Higher levels work on records and files (of records)

File: collection of pages

- insert/delete/modify record

- get(record\_id) a record

- scan all records

Page: collection of records

- typically *fixed* size (8kb in PostgreSQL)

May be stored in multiple OS files spanning multiple disks



# Units that we'll care about

Ignore CPU cost

Ignore RAM cost

**B** # *data* pages on disk for relation

**R** # records per data page

**D** avg time to read/write data page to/from disk

Simplifies life when computing costs

Very rough approximation, but OK for now

ignores prefetching, bulk writes/reads, CPU/RAM

# Unordered Heap Files

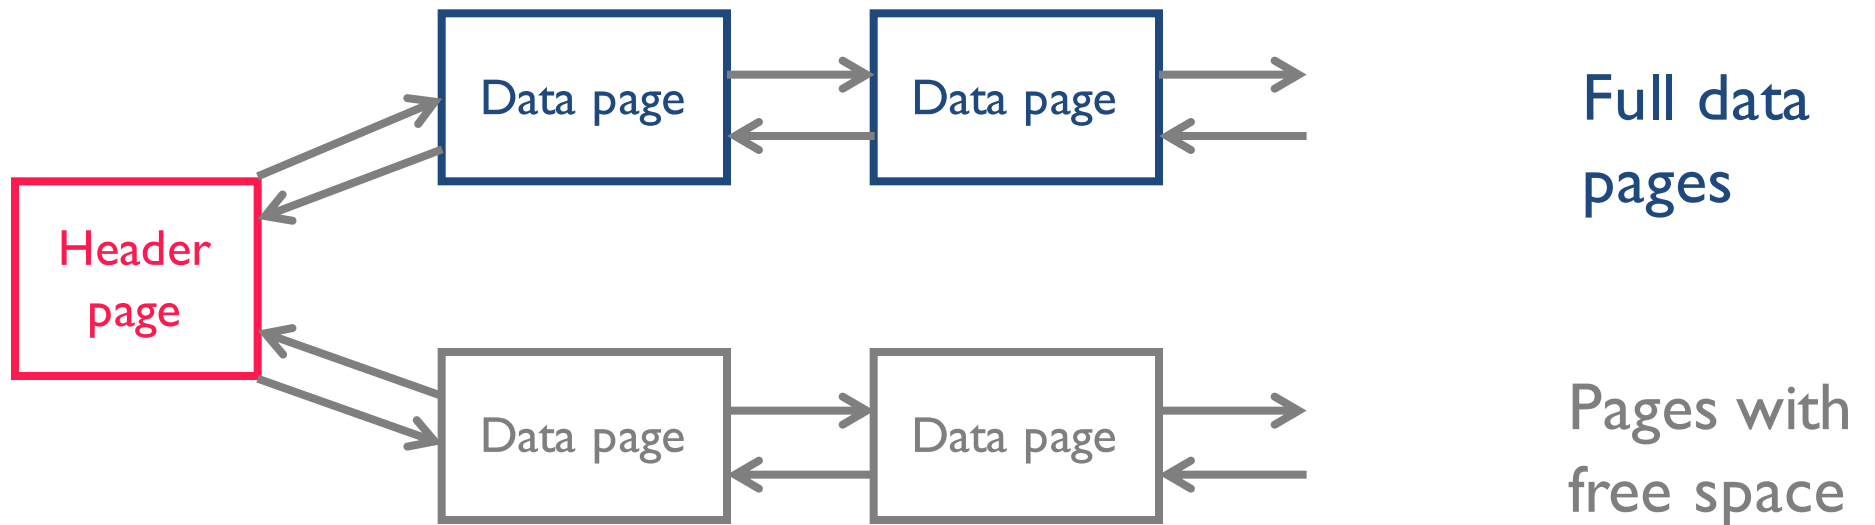
Collection of records (no order)

As add/rm records, pages de/allocated

To support record level ops, need to track:

- pages in file
- free space on pages
- records on page


# Heap File



Header page and heap file pointers stored in catalog

Data page = 2 pointers + data

# Use a directory



Directory entries track #free bytes on data pages

Directory is collection of pages

# Indexes

Heap files can get data  
by rid  
by sequential scan


Queries use *qualifications* (predicates)  
find students in “CS”  
find students from CA

## Indexes


file structures for value-based queries  
B+-tree index (~1970s)  
Hash index

Overview! Details in 4112


# B+ Tree




Query: **SELECT \* WHERE val = 14**




# B+ Tree




Query: **SELECT val WHERE val = 14**  
(index only)



# B+ Tree




Query: **SELECT \* WHERE val = 55**






# B+ Tree



# B+ Tree



Query: **SELECT \* WHERE val > 20**

# Some numbers (8kb pages)

How many levels?

fill-factor: ~66%

~300 entries per directory page

if fill completely

height 2:  $300^3 \sim 27$  Million

height 3:  $300^4 \sim 8.1$  Billion

Top levels often in memory

Level 3 only 90k pages ~750MB

# Hash Index

1


$$h(v) = v \% 3$$

0, 6

4, 13

5, 14


# Hash Index




# Hash Index

11


$$h(v) = v \% 3$$



# Hash Index



# Hash Index



Good for equality selections

Index = data pages + overflow data pages

Hash function  $h(v)$  takes as input the *search key*



# Costs

Three file types

Heap, B+ Tree, Hash

Operations we care about

Scan all data     `SELECT * FROM R`

Equality            `SELECT * FROM R WHERE x = I`

Range              `SELECT * FROM R WHERE x > 10 and x < 50`

Insert record

Delete record

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything				
Equality				
Range				
Insert				
Delete				

**B** # data pages

**D** time to read/write page

**M** # pages in range query

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD			
Equality	0.5BD			
Range	BD			
Insert	2D			
Delete	Search + D			

## Heap File

equality on a key. How many results?

**B** # data pages

**D** time to read/write page

**M** # pages in range query

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD	BD		
Equality	0.5BD	$D(\log_2 B)$		
Range	BD	$D(\log_2 B + M)$		
Insert	2D	Search + BD		
Delete	Search + D	Search + BD		

## Heap File

equality on a key. How many results?

## Sorted File

files compacted after deletion

**B** # data pages

**D** time to read/write page

**M** # pages in range query

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD	BD	1.2BD	
Equality	0.5BD	$D(\log_2 B)$	$D(\log_{80} B + 1)$	
Range	BD	$D(\log_2 B + M)$	$D(\log_{80} B + M)$	
Insert	2D	Search + BD	$D(\log_{80} B)$	
Delete	Search + D	Search + BD	$D(\log_{80} B)$	

## Heap File

equality on a key. How many results?

## Sorted File

files compacted after deletion

## B+ Tree

100 entries/directory page

80% fill factor

**B** # data pages

**D** time to read/write page

**M** # pages in range query

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD	BD	1.2BD	1.2BD
Equality	0.5BD	$D(\log_2 B)$	$D(\log_{80} B + 1)$	D
Range	BD	$D(\log_2 B + M)$	$D(\log_{80} B + M)$	1.2BD
Insert	2D	Search + BD	$D(\log_{80} B)$	2D
Delete	Search + D	Search + BD	$D(\log_{80} B)$	2D

## Heap File

equality on a key. How many results?

## Sorted File

files compacted after deletion

## B+ Tree

100 entries/directory page

80% fill factor

## Hash index

no overflow

80% fill factor

**B** # data pages

**D** time to read/write page

**M** # pages in range query

# How to pick?

Depends on your queries (workload)

Which relations?

Which attributes?

Which types of predicates ( $=$ ,  $<$ ,  $>$ )

*Selectivity*

Insert/delete/update queries? how many?

# How to choose indexes?

## Considerations

which relations should have indexes?

on what attributes?

how many indexes?

what type of index (hash/tree)?



# Naïve Algorithm

for each query in order of importance  
    calculate best cost using current indexes  
    if there's best plan with a new index  
        create it

Why not create every index?

    update queries slowed down (upkeep costs)  
    takes up space

# High level guidelines

Check the WHERE clauses

- attributes in WHERE are search/index keys

- equality predicate → hash index

- range predicate → tree index

Multi-attribute search keys supported

- order of attributes matters for range queries

- may enable queries that don't look at data pages (*index-only*)

# Summary

Design depends on economics, access cost ratios

Disk still dominant wrt cost/capacity ratio

Many physical layouts for files

same APIs, difference performance

remember physical independence

## Indexes

Structures to speed up read queries

Multiple indexes possible

Decision depends on workload