

## Administrivia

Project 2 not ready :(

How did the evaluations go?

How did the projects go?

# L9

## Query Execution & Optimization

Eugene Wu  
Fall 2015

Remember in the history of SQL, the CODASYL folks were making the performance argument, that there's no way that SQL could run fast? Well history has proved them wrong, and we'll talk enough to give you a feel for how we went about making this fast

# Steps for a New Application

## Requirements

what are you going to build?

## Conceptual Database Design

pen-and-pencil description

## Logical Design

formal database schema

## Schema Refinement:

fix potential problems, normalization

## Physical Database Design

optimize for speed/storage

Optimization

## App/Security Design

prevent security problems

# Recall

## Relational algebra

equivalence: multiple stmts for same query  
some statements (much) faster than others

### Which is faster?

- a.  $\sigma_{v=1}(R \times T)$
- b.  $\sigma_{v=1}(\sigma_{v=1}(R) \times T)$

$|R| = |T| = 10$  pages. 100? 1M?

# unique values in R = 1. 100? 1M?  $\longleftarrow$  selectivity!

In the first relational algebra query, we first perform a cross product, which will return  $R \times T$  pages, then we filter by  $v=1$

In the second, we first filter R, and then perform the cross product

in the extreme case, where there are no R records with  $v = 1$ , then this is an obviously better option.

$2 \times (R \times T) + R \times T$ : for each R tuple, read each T and generate result ( $R \times T$ ). write it out. then read it back in for filter

vs

$(1 + \text{selectivity}) \times R + (\text{selectivity} \times R) \times T$

# Overview of Query Optimization

SQL → query plan

How plans are executed

Some implementations of operators

Cost estimation of a plan

Selectivity

System R dynamic programming

All ideas from System R's "Selinger Optimizer" 1979

First, we need to talk about how exactly to translate SQL into a standardized data structure that can be manipulated, compared, optimized. That form is the query plan

Go through the bottom up execution of these query plans

Look at some alternative operator implementation – different join algorithms, etc

Because of this there is a distinction between a logical operator such as the greek symbols we have seen

(describe the result semantics)

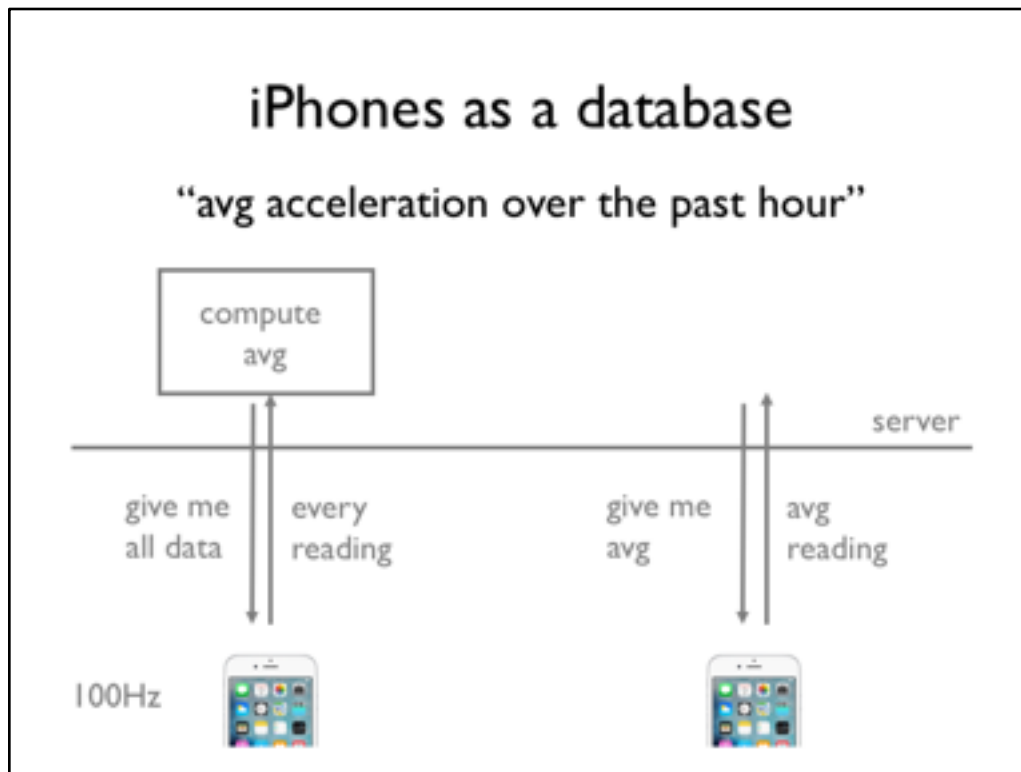
and the physical implementation of the operator or physical operator, that is exactly how it should be executed

pipelining

access paths

Optimizers are hugely important – often the trade secret of a database. not surprising to have hundreds of programmer years in an optimizer like Oracle or IBM's

A huge chunk of database research, at the core, are optimization algorithms that could be added to an optimizer



100hz = 100 samples per second, 6000 per minute, 360,000 per hour  
example where doing the avg computation on the phone (pushing it down) makes sense.

## SQL $\rightarrow$ Query Plan

SELECT a FROM R	$\pi_a(R)$	$\begin{array}{c} \pi_a \\   \\ R \end{array}$
SELECT a FROM R WHERE a > 10	$\pi_a(\sigma_{a>10}(R))$	$\begin{array}{c} \pi_a \\   \\ \sigma_{a>10} \\   \\ R \end{array}$
SELECT a FROM R JOIN S ON R.b = S.b	$\pi_a(\bowtie_b(R,S))$	$\begin{array}{c} \pi_a \\   \\ \bowtie_b \\ / \quad \backslash \\ R \quad S \end{array}$

arguments are children

# Query Evaluation

Push vs Pull?

Push

Operators are input-driven

As operator (say reading input table) gets data, push it to parent operator.

Pull

Operators are demand-driven

If parent says "give me next result", then do the work

Are cursors push or pull?



# Query Evaluation

Naïve execution (operator at a time)

read R

filter  $a > 10$  and write out

read and project a

Cost:  $B + M + M$

SELECT a  
FROM R  
WHERE  $a > 10$

$\pi_a$   
|  
 $\sigma_{a > 10}$   
|  
R

**B** # data pages

**M** # pages matched in  
WHERE clause

Could we do better?

I read R, maybe after the filter the result doesn't fit into memory, or just barely, so I need to write it out. then I go to next operator, and read its data to execute the projection

# Query Evaluation

Pipelined exec (tuple/page at a time)

read first page of R, pass to  $\sigma$

filter  $a > 10$  and pass to  $\pi$

project a

(all operators run concurrently)

Cost: B

SELECT a  
FROM R  
WHERE  $a > 10$

$\pi_a$   
|  
 $\sigma_{a>10}$   
|  
R

**B** # data pages

**M** # pages matched in  
WHERE clause

Note: can't pipeline some operators!

e.g., sort, some joins, aggregates

why?

Not exactly correct, but provides the intuition of why pipelining is a good idea. In reality, each operator is often implemented using an iterator interface with `get_next()` calls.

The user would call `next()` on the root node, `pi` in this case, and the query execution will do just the work to compute the next result tuple.

In this case, if I only called `get_next()` once, we would only need to read a single page!

Keep in mind that this is an example of pipelined execution – all of the operators are running at the same time on the same or different pieces of data. If each operator were a separate CPU, then they are running at the same time and don't need to wait on each other.

Remember the monotonicity property from talking about relational operators? there are operators such as some join operators that are blocking, meaning

Could we do better?

# Query Evaluation

What if R is indexed?

Hash index

Not appropriate

B+Tree index

use  $a > 10$  to find initial data page

scan leaf data pages

Cost:  $\log_f B + M$

SELECT a  
FROM R  
WHERE  $a > 10$

$\pi_a$   
|  
 $\sigma_{a > 10}$   
|  
R

**B** # data pages

**M** # pages matched in  
WHERE clause

Not exactly correct, but provides the intuition of why pipelining is a good idea. In reality, each operator is often implemented using an iterator interface with `get_next()` calls.

The user would call `next()` on the root node, `pi` in this case, and the query execution will do just the work to compute the next result tuple.

In this case, if I only called `get_next()` once, we would only need to read a single page!

Could we do better?

## Access Paths

Choice of how to access input data is called the  
**Access Path**

file scan or

index + matching condition (e.g.,  $a > 10$ )

# Access Paths

## Sequential Scan

doesn't accept any matching conditions

## Hash index search key $\langle a, b, c \rangle$

accepts conjunction of equality conditions on *all* search keys

e.g.,  $a = 1$  and  $b = 5$  and  $c = 5$

will  $(a = 1$  and  $b = 5)$  work? why?

## Tree index search key $\langle a, b, c \rangle$

accepts conjunction of terms of *prefix* of search keys

e.g.,  $a > 1$  and  $b = 5$  and  $c < 5$

will  $(a > 1$  and  $b = 5)$  work?

will  $(a > 1$  and  $c > 9)$  work?

$(a > 1$  and  $c > 9)$  will work because  $a > 1$  uses  $a$ , which is a prefix of the search key. When we get to the left-most leaf data node that matches  $a > 1$ , we will then scan towards the left  
if the data has been sorted at the leaves, then it's just a sequential scan for all pages that contain tuples that match  $a > 1$

## How to pick Access Paths?

### Selectivity

ratio of # outputs satisfying predicates vs # inputs

0.01 means 1 output tuple for every 100 input tuples

### Assume

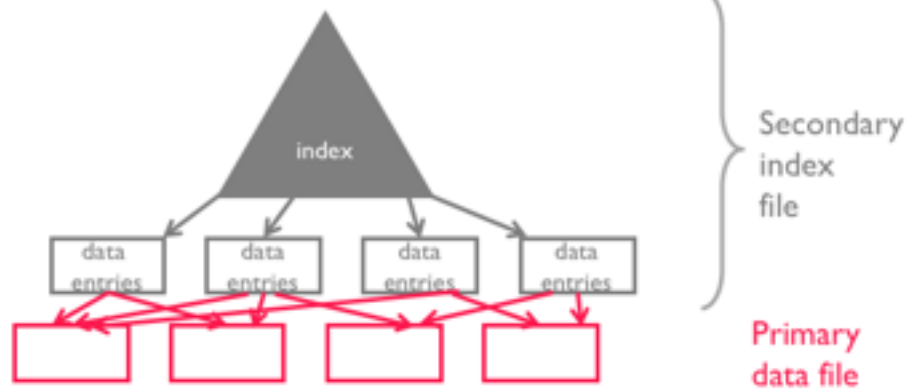
attribute selectivity is independent

if  $\text{selectivity}(a=1) = 0.1$ ,  $\text{selectivity}(b>3) = 0.6$

$\text{selectivity}(a=1 \text{ and } b>3) = 0.1 * 0.6 = 0.06$

Why does selectivity matter?

## High level index structure



What is a data entry?

actual data record

<search key value, rid>

<search key value, rid\_list>

baseline is sequential scan

if storing data in index as RID, then one page access per record

same with hash tables

# How to pick Access Paths?

Hash index on  $\langle a, b, c \rangle$

$a = 1, b = 1, c = 1$  how to estimate selectivity?

1. pre-compute attribute statistics by scanning data  
e.g., a has 100 values, b has 200 values, c has 1 value  
selectivity =  $1 / (100 * 200 * 1)$
2. How many distinct values does hash index have?  
e.g., 1000 distinct values in hash index
3. make a number up  
"default estimate" is the fancy term

- 1) look at each attribute individually
- 2) look at combination of all attributes and their distinctness
- 3) make something up



## System Catalog Keeps Statistics

### System R

NCARD	"relation cardinality" # tuples in relation
TCARD	# pages relation occupies
ICARD	# keys (distinct values) in index
NINDEX	pages occupied by index
	min and max keys in indexes

Statistics were expensive in 1979!

Super elegant: catalog stored in relations too!

What is a statistic? It's a data structure (sometimes just a single number) that describes enough about the data to estimate the selectivities – and thus the costs for example, you could get perfect information by just running the query. so it's a trade off between accuracy and time

for example, the index is a compressed, summarized version of the table for ONLY the search key attributes,  
but presumably those are the attribute you CARE about estimating selectivities correctly

if stats stored in database, then estimates are just queries.

## What Optimization Options Do We Have?

Access Path ✓

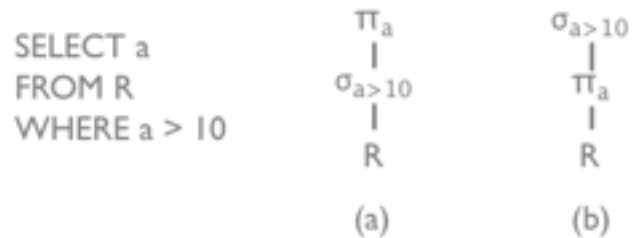
Predicate push-down

Join implementation

Join ordering

In general, depends on operator implementations. So let's take a look

## Predicate Push Down



Which is faster if B+ Tree index: (a) or (b)?

(a)  $\log_f(B) + M$  pages

(b) B pages

It's a Good Idea, especially when we look at Joins

do (a) and (b) do the same things?

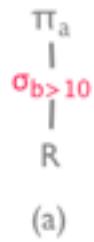
if you're thinking, well it's obvious that a and b are equivalent plans, then you're right!

in (b), the selection predicate  $a > 10$  is at the top, and in (a), it is "pushed down" lower into the query plan

What would be a rule we could create that tells us to pick the faster one?

## Predicate Push Down

SELECT a  
FROM R  
WHERE a > 10



You have to be a bit careful about how you do the predicate pushdown and when you are allowed to

## Projection with DISTINCT clause

need to deduplicate e.g.,  $\pi_{\text{rating}} \text{Sailors}$

### Two basic approaches

Sort: fundamental database operation

sort on rating, remove dups on scan of sorted data

Hash:

partition into N buckets

sort each bucket and remove dups

### Index on projected fields

scan the index pages, avoid reading data

Sorting is

# The Join

**Core database operation**

join of 100 tables common in enterprise apps

**Join algorithms is a large area of research**

e.g., distributed, temporal, geographic, multi-dim, range,  
sensors, graphs, etc

**Discuss three basic joins**

nested loops, indexed nested loops, hash join

**Best join implementation depends on the query, the data,  
the indices, hardware, etc**

if you squint, almost everything can be viewed as a join, and its ideas are constantly being rediscovered

For example, graph analysis – we saw that it is logically a join

When you have a dataset of actors and download a dataset of movies to analyze them, that's a join in space – performing type of join

When you high five a friend, that's a join in space.. and time

When a florescent marker binds to a target protein, that's a chemical join

By thinking of this way, we can make use of everything we know about joins

One could say it's a fundamental operation in life

## Datasets

```
from collections import defaultdict
from random import randint

M = 10000
N = 1000
outer = [ [randint(0, 1000), randint(0, 1000)]
           for i in xrange(M)]
inner = [ [randint(0, 1000), randint(0, 1000)]
          for i in xrange(N)]

# outer ⋈ inner
# outer JOIN inner ON outer.1 = inner.1
```

OK, we'll illustrate a few join algorithms using python code. You should be able to copy and past this code to run it

As opposed to cost estimation before, we'll be pretty general about join costs. Partly because you'll go into substantially more detail in 4112, I want you to understand nested loops and indexed nested loops join, for the other joins, just the properties

## Nested Loops Join

```
for row in outer:  
    for irow in inner:  
        if row[0] == irow[0]:    # could be any check  
            yield (row, irow)
```

### Very flexible

Equality check can be replaced with any condition

Incremental algorithm

Cost:  $M + MN$

Is this the same as a cross product?

Different than cross product because output size is not  $M*N$   
but otherwise, the algorithm is the same!



## Indexed Nested Loops Join

```
for row in outer:  
    for irow in index.get(row[0], []):  
        yield (row, irow)
```

### Slightly less flexible

Only supports conditions that the index supports

Assume hash index, 50 tuples per page, equality join  
on sid, 5% of outer tuples match

Cost:  $M + (50M * 0.05 * I_{\text{page}})$

You could imagine, if the inner table is small enough, building the hash table on the fly and running indexed nested loops join  
or, building hash tables on both outer and inner tables  
There are a bunch of variations of this idea, and they primarily think about how to best use RAM vs going to disk, as well as using sequential access of the disk

## Sort Merge Join

Sort outer and inner tables on join key

Cost: 2-3 scans of each table

Merge the tables and compute the join

Cost: 1 scan of each table

Overall Properties

cost:  $3(M+N)$  to  $4(M+N)$

results are sorted

highly sequential access

(weapon of choice for very large datasets)

mention mcsherry's blog post – often for huge datasets, sort merge can out perform hash lookups due to random access vs sequential

## Where we are

We've discussed

- Optimizations for a single operator

- Different types of access paths, join operators

- Simple optimizations e.g., predicate push-down

What about for multiple operators?

- System R Optimizer

# System R Optimizer

Granddaddy of all existing optimizers

don't go for best plan, go for *least worst plan*

- **Cost Estimator**

"predict" cost of query from statistics

Includes CPU, disk, memory, etc (can get sophisticated!)

It's an art

- **Plan Space**

avoid cross product

push selections & projections to leaves as much as possible

only join ordering remaining

# Cost Estimation

`estimate(operator, inputs, stats) → cost`

estimate cost for each operator

depends on input *cardinalities* (# tuples)

discussed earlier in lecture

estimate output size for each operator

used by `estimate()` of parent operator

use selectivity. assume attributes are independent

Try it in PostgreSQL: `EXPLAIN <query>;`

you'll note that I don't have any unit for the cost. That's because estimators tend to be so wildly off that they bear almost no relation to wall clock time.

If you

## Estimate Size of Output

```
SELECT      *  
FROM        R1, ..., Rn  
WHERE       term1 AND ... AND termm
```

Query input size

$$|R1| * \dots * |Rn|$$

Term selectivity

$$\text{col} = v \quad 1 / \text{ICARD}_{\text{col}}$$

$$\text{col1} = \text{col2} \quad 1 / \max(\text{ICARD}_{\text{col1}}, \text{ICARD}_{\text{col2}})$$

$$\text{col} > v \quad (\max_{\text{col}} - v) / (\max_{\text{col}} - \min_{\text{col}})$$

Query output size

$$|R1| * \dots * |Rn| * \text{term}_1 \text{selectivity} * \dots * \text{term}_m \text{selectivity}$$

## Try it out

R.sid = S.sid selectivity 0.01

R.bid selectivity 0.05

$|R| = M$

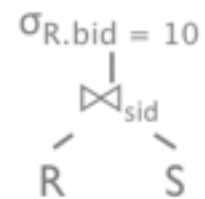
$|S| = N$

Cost:  $M + MN$

projection is pipelined

# outputs:  $0.0005MN$

```
SELECT *  
FROM R, S  
WHERE R.sid = S.sid  
AND R.bid = 10
```



## Try it out

R.sid = S.sid selectivity 0.01

R.bid selectivity 0.05

$|R| = M$

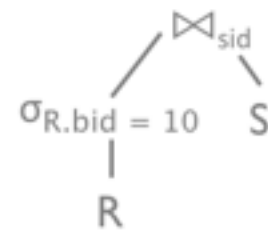
$|S| = N$

Cost:  $M + (0.05MN)$

projection is pipelined

# outputs:  $0.0005MN$

```
SELECT *  
FROM R, S  
WHERE R.sid = S.sid  
AND R.bid = 10
```





## Plan Space

$A \bowtie B \bowtie C$



How many  
plans?

(AB)C  
A(BC)  
(AC)B

(BC)A  
B(AC)  
(CA)B

(BA)C  
B(AC)  
(CB)A

A(CB)  
C(AB)  
C(BA)

# parenthetizations \* #strings

## Plan Space

# parenthetizations \* #strings

A: (A)

AB: (AB)

ABC: ((AB)C), (A(BC))

ABCD: (((AB)C)D), ((A(BC))D), ((AB)(CD)), (A((BC)D)), (A(B(CD)))

paren(n) choose(2(N-1), (N-1)) / N

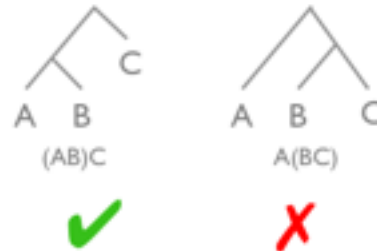
choose(2(N-1), (N-1)) / N \* N!

N=1, #plans = 17,643,225,600

## Reducing the Plan Space

$\text{parens}(N) = 1$

Only left-deep plans  
ensures pipelining



### Dynamic Programming

Idea: If considering  $((ABC)DE)$   
compute best  $(ABC)$ , cache, and reuse  
figure out best way to combine with  $(DE)$

### Dynamic Programming Algorithm

compute best join size 1, then size 2, ...  
 $\sim O(2^N)$

pipelining: right deep tree on the right. before A can join with it's inner relation (result of B join C), need to wait until B join C is completed

## Summary

Query optimization is a deep, complex topic

Pipelined plan execution

Different types of joins

Cost estimation of single and multiple operators

Join ordering is hard

# Reducing the Plan Space

Dynamic Programming Algorithm

compute best join size 1, then size 2, ...

R = relations to join

for n in {1,... |R|}

for S in {all size n subsets of R}

bestjoin(S) = A join S-A

where A is single relation that minimizes:

cost(bestjoin(S-A)) +

min cost to join A to (S-A) +

min access cost for A

## Selinger Algorithm $N = 1$

bestjoin(ABC), only nested loops join

$N = 1$

A = best way to access A

B = best way to access B

C = best way to access C

cost: N relations

## Selinger Algorithm $N = 2$

bestjoin(ABC)

$N = 2$

A,B = AB or BA

A,C = AC or CA

B,C = BC or CB

cost:  $\text{choose}(N, 2) * 2$

## Selinger Algorithm $N = 3$

bestjoin(ABC)

$N = 3$

A,B,C = (AB)C or (AC)B or (BC)A

cost:  $\text{choose}(N, 3) * 3$



## Selinger Algorithm Cost

cost = # subsets \* # options per subset  
set of relations R,  $N = |R|$

#subsets = choose( $N, 1$ ) + choose( $N, 2$ ) + choose( $N, 3$ )...  
=  $2^N$

#options =  $k < N$  ways to remove a relation A +  
2 ways to join A with R-A  
 $< 2N$

Cost =  $2N * 2^N$

2 ways to join A with R-A only because considering NL join only. otherwise it should be  $2 * M$ , where M is the number of possible join algorithms