

L9

Query Execution & Optimization

Eugene Wu
Fall 2015

Remember in the history of SQL, the CODASYL folks were making the performance argument, that there's no way that SQL could run fast? Well history has proved them wrong, and we'll talk enough to give you a feel for how we went about making this fast

Steps for a New Application

Requirements

what are you going to build?

Conceptual Database Design

pen-and-pencil description

Logical Design

formal database schema

Schema Refinement:

fix potential problems, normalization

Physical Database Design

optimize for speed/storage

Optimization

App/Security Design

prevent security problems

Recall

Relational algebra

equivalence: multiple stmts for same query
some statements (much) faster than others

Which is faster?

- a. $\sigma_{v=1}(R \times T)$
- b. $\sigma_{v=1}(\sigma_{v=1}(R) \times T)$

$|R| = |T| = 10$ pages. 100? 1M?

unique values in R = 1. 100? 1M? \longleftarrow selectivity!

In the first relational algebra query, we first perform a cross product, which will return $R \times T$ pages, then we filter by $v=1$

In the second, we first filter R, and then perform the cross product

in the extreme case, where there are no R records with $v = 1$, then this is an obviously better option.

$2 \times (R \times T) + R \times T$: for each R tuple, read each T and generate result ($R \times T$). write it out. then read it back in for filter

vs

$(1 + \text{selectivity}) \times R + (\text{selectivity} \times R) \times T$

Overview of Query Optimization

SQL → query plan

How plans are executed

Some implementations of operators

Cost estimation of a plan

Selectivity

System R dynamic programming

All ideas from System R's "Selinger Optimizer" 1979

First, we need to talk about how exactly to translate SQL into a standardized data structure that can be manipulated, compared, optimized. That form is the query plan

Go through the bottom up execution of these query plans

Look at some alternative operator implementation – different join algorithms, etc

Because of this there is a distinction between a logical operator such as the greek symbols we have seen

(describe the result semantics)

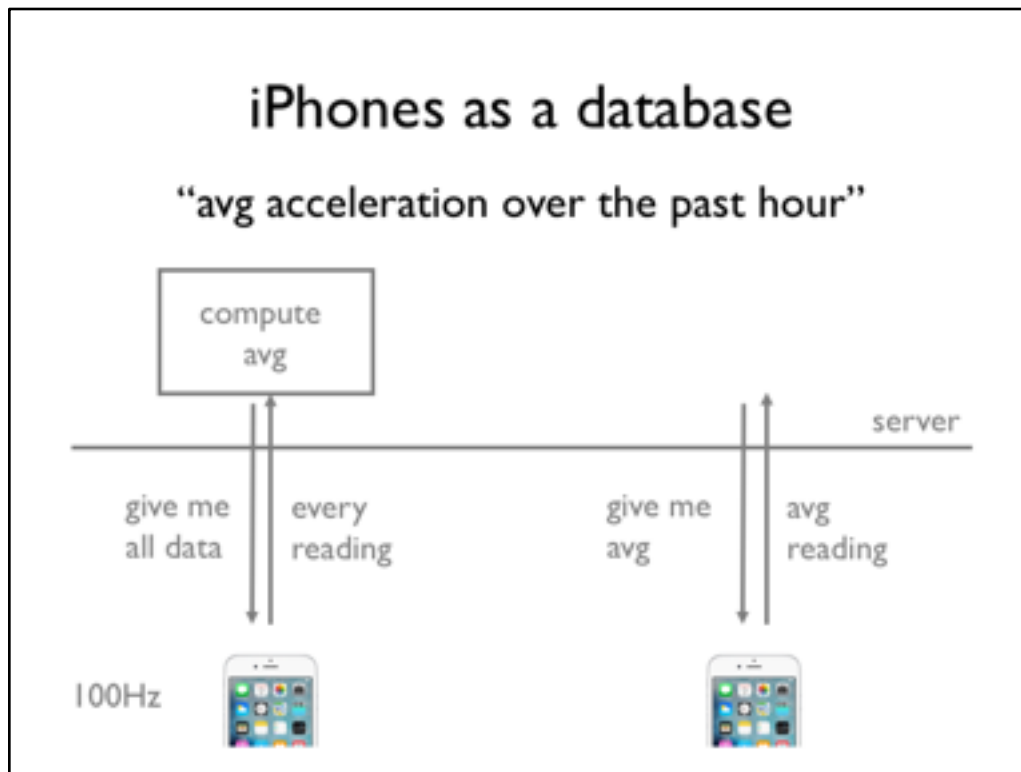
and the physical implementation of the operator or physical operator, that is exactly how it should be executed

pipelining

access paths

Optimizers are hugely important – often the trade secret of a database. not surprising to have hundreds of programmer years in an optimizer like Oracle or IBM's

A huge chunk of database research, at the core, are optimization algorithms that could be added to an optimizer



100hz = 100 samples per second, 6000 per minute, 360,000 per hour
example where doing the avg computation on the phone (pushing it down) makes sense.

SQL \rightarrow Query Plan

SELECT a FROM R	$\pi_a(R)$	$\begin{array}{c} \pi_a \\ \\ R \end{array}$
SELECT a FROM R WHERE a > 10	$\pi_a(\sigma_{a>10}(R))$	$\begin{array}{c} \pi_a \\ \\ \sigma_{a>10} \\ \\ R \end{array}$
SELECT a FROM R JOIN S ON R.b = S.b	$\pi_a(\bowtie_b(R,S))$	$\begin{array}{c} \pi_a \\ \\ \bowtie_b \\ / \quad \backslash \\ R \quad S \end{array}$

arguments are children

Query Evaluation

Push vs Pull?

Push

Operators are input-driven

As operator (say reading input table) gets data, push it to parent operator.

Pull

Operators are demand-driven

If parent says "give me next result", then do the work

Are cursors push or pull?

Query Evaluation

Naïve execution (operator at a time)

read R

filter $a > 10$ and write out

read and project a

Cost: $B + M + M$

SELECT a
FROM R
WHERE $a > 10$

π_a
|
 $\sigma_{a > 10}$
|
R

B # data pages

M # pages matched in
WHERE clause

Could we do better?

I read R, maybe after the filter the result doesn't fit into memory, or just barely, so I need to write it out. then I go to next operator, and read its data to execute the projection

Query Evaluation

Pipelined exec (tuple/page at a time)

read first page of R, pass to σ

filter $a > 10$ and pass to π

project a

(all operators run concurrently)

Cost: B

SELECT a
FROM R
WHERE a > 10

π_a
|
 $\sigma_{a>10}$
|
R

B # data pages

M # pages matched in
WHERE clause

Note: can't pipeline some operators!

e.g., sort, some joins, aggregates

why?

Not exactly correct, but provides the intuition of why pipelining is a good idea. In reality, each operator is often implemented using an iterator interface with `get_next()` calls.

The user would call `next()` on the root node, `pi` in this case, and the query execution will do just the work to compute the next result tuple.

In this case, if I only called `get_next()` once, we would only need to read a single page!

Keep in mind that this is an example of pipelined execution – all of the operators are running at the same time on the same or different pieces of data. If each operator were a separate CPU, then they are running at the same time and don't need to wait on each other.

Remember the monotonicity property from talking about relational operators? there are operators such as some join operators that are blocking, meaning

Could we do better?

Query Evaluation

What if R is indexed?

Hash index

Not appropriate

B+Tree index

use $a > 10$ to find initial data page

scan leaf data pages

Cost: $\log_f B + M$

SELECT a
FROM R
WHERE $a > 10$

π_a
|
 $\sigma_{a > 10}$
|
R

B # data pages

M # pages matched in
WHERE clause

Not exactly correct, but provides the intuition of why pipelining is a good idea.
In reality, each operator is often implemented using an iterator interface with `get_next()` calls.

The user would call `next()` on the root node, `pi` in this case, and the query execution will do just the work to compute the next result tuple.

In this case, if I only called `get_next()` once, we would only need to read a single page!

Could we do better?

Access Paths

Choice of how to access input data is called the
Access Path

file scan or

index + matching condition (e.g., $a > 10$)

Access Paths

Sequential Scan

doesn't accept any matching conditions

Hash index search key $\langle a, b, c \rangle$

accepts conjunction of equality conditions on *all* search keys

e.g., $a = 1$ and $b = 5$ and $c = 5$

will $(a = 1$ and $b = 5)$ work? why?

Tree index search key $\langle a, b, c \rangle$

accepts conjunction of terms of *prefix* of search keys

e.g., $a > 1$ and $b = 5$ and $c < 5$

will $(a > 1$ and $b = 5)$ work?

will $(a > 1$ and $c > 9)$ work?

$(a > 1$ and $c > 9)$ will work because $a > 1$ uses a , which is a prefix of the search key. When we get to the left-most leaf data node that matches $a > 1$, we will then scan towards the left
if the data has been sorted at the leaves, then it's just a sequential scan for all pages that contain tuples that match $a > 1$

How to pick Access Paths?

Selectivity

ratio of # outputs satisfying predicates vs # inputs

0.01 means 1 output tuple for every 100 input tuples

Assume


attribute selectivity is independent

if $\text{selectivity}(a=1) = 0.1$, $\text{selectivity}(b>3) = 0.6$

$\text{selectivity}(a=1 \text{ and } b>3) = 0.1 * 0.6 = 0.06$

Why does selectivity matter?

High level index structure



What is a data entry?

actual data record

<search key value, rid>

<search key value, rid_list>

baseline is sequential scan

if storing data in index as RID, then one page access per record

same with hash tables

How to pick Access Paths?

Hash index on $\langle a, b, c \rangle$

$a = 1, b = 1, c = 1$ how to estimate selectivity?

1. pre-compute attribute statistics by scanning data
e.g., a has 100 values, b has 200 values, c has 1 value
selectivity = $1 / (100 * 200 * 1)$
2. How many distinct values does hash index have?
e.g., 1000 distinct values in hash index
3. make a number up
"default estimate" is the fancy term

- 1) look at each attribute individually
- 2) look at combination of all attributes and their distinctness
- 3) make something up

System Catalog Keeps Statistics

System R

NCARD	"relation cardinality" # tuples in relation
TCARD	# pages relation occupies
ICARD	# keys (distinct values) in index
NINDEX	pages occupied by index
	min and max keys in indexes

Statistics were expensive in 1979!

Super elegant: catalog stored in relations too!

What is a statistic? It's a data structure (sometimes just a single number) that describes enough about the data to estimate the selectivities – and thus the costs for example, you could get perfect information by just running the query. so it's a trade off between accuracy and time

for example, the index is a compressed, summarized version of the table for ONLY the search key attributes,
but presumably those are the attribute you CARE about estimating selectivities correctly

if stats stored in database, then estimates are just queries.

What Optimization Options Do We Have?

Access Path ✓


Predicate push-down

Join implementation

Join ordering

In general, depends on operator implementations. So let's take a look

Predicate Push Down



Which is faster if B+ Tree index: (a) or (b)?

(a) $\log_f(B) + M$ pages

(b) B pages

It's a Good Idea, especially when we look at Joins

do (a) and (b) do the same things?

if you're thinking, well it's obvious that a and b are equivalent plans, then you're right!

in (b), the selection predicate $a > 10$ is at the top, and in (a), it is "pushed down" lower into the query plan

What would be a rule we could create that tells us to pick the faster one?

Projection with DISTINCT clause

need to deduplicate e.g., π_{rating} Sailors

Two basic approaches

Sort: fundamental database operation

sort on rating, remove dups on scan of sorted data

Hash:

partition into N buckets

sort each bucket and remove dups

Index on projected fields

scan the index pages, avoid reading data

Sorting is

The Join

Core database operation

join of 100 tables common in enterprise apps

Join algorithms is a large area of research

e.g., distributed, temporal, geographic, multi-dim, range, sensors, graphs, etc

Discuss three basic joins

nested loops, indexed nested loops, hash join

Best join implementation depends on the query, the data, the indices, hardware, etc

if you squint, almost everything can be viewed as a join, and its ideas are constantly being rediscovered

For example, graph analysis – we saw that it is logically a join

When you have a dataset of actors and download a dataset of movies to analyze them, that's a join in space – performing type of join

When you high five a friend, that's a join in space.. and time

When a florescent marker binds to a target protein, that's a chemical join

By thinking of this way, we can make use of everything we know about joins

One could say it's a fundamental operation in life

Datasets

```
from collections import defaultdict
from random import randint

M = 10000
N = 1000
outer = [ [randint(0, 1000), randint(0, 1000)]
           for i in xrange(M)]
inner = [ [randint(0, 1000), randint(0, 1000)]
          for i in xrange(N)]

# outer ⋈ inner
# outer JOIN inner ON outer.1 = inner.1
```

OK, we'll illustrate a few join algorithms using python code. You should be able to copy and past this code to run it

As opposed to cost estimation before, we'll be pretty general about join costs. Partly because you'll go into substantially more detail in 4112, I want you to understand nested loops and indexed nested loops join, for the other joins, just the properties

Nested Loops Join

```
for row in outer:  
    for irow in inner:  
        if row[0] == irow[0]:    # could be any check  
            yield (row, irow)
```

Very flexible

Equality check can be replaced with any condition

Incremental algorithm

Cost: $M + MN$

Is this the same as a cross product?

Different than cross product because output size is not $M*N$
but otherwise, the algorithm is the same!

Indexed Nested Loops Join

```
for row in outer:  
    for irow in index.get(row[0], []):  
        yield (row, irow)
```

Slightly less flexible

Only supports conditions that the index supports

Assume hash index, 50 tuples per page, equality join
on sid, 5% of outer tuples match

Cost: $M + (50M * 0.05 * I_{\text{page}})$

$M \text{ pages} * 50 \text{ tuples per page} = 50M \text{ tuples in outer}$. Each is a probe of the index, and 5% of the probes are a hit that require accessing a data page (1 page)
+ M pages to read for the outer table.

You could imagine, if the inner table is small enough, building the hash table on the fly and running indexed nested loops join
or, building hash tables on both outer and inner tables
There are a bunch of variations of this idea, and they primarily think about how to best use RAM vs going to disk, as well as using sequential access of the disk

Sort Merge Join

Sort outer and inner tables on join key

Cost: 2-3 scans of each table

Merge the tables and compute the join

Cost: 1 scan of each table

Overall Properties

cost: $3(M+N)$ to $4(M+N)$

results are sorted

highly sequential access

(weapon of choice for very large datasets)

mention mcsherry's blog post – often for huge datasets, sort merge can out perform hash lookups due to random access vs sequential

Quick Recap

Single operator optimizations

- Access paths
- Primary vs secondary index costs
- Projection/distinct
- Predicate/project push downs

2 operators aka Joins

- Nested loops, index nested loops, sort merge

Selectivity estimation

- Statistics and simple models

Using an Index for Selections

Cost depends on #qualifying tuples.

- Cost of finding qualifying data entries (typically small) plus cost of retrieving records (could be large).

- In example, assuming uniform distribution of names, about 10% of tuples qualify (100 pages, 10000 tuples). With a “clustered” index, cost is little more than 100 I/Os; if “unclustered,” up to 10000 I/Os!

Where we are

We've discussed

- Optimizations for a single operator

- Different types of access paths, join operators

- Simple optimizations e.g., predicate push-down

What about for multiple operators?

- System R Optimizer

Selinger Optimizer

Granddaddy of all existing optimizers

don't go for best plan, go for *least worst plan*

2 Big Ideas

1. **Cost Estimator**

"predict" cost of query from statistics

Includes CPU, disk, memory, etc (can get sophisticated!)

It's an art

2. **Plan Space**

avoid cross product

push selections & projections to leaves as much as possible

only join ordering remaining

Selinger Optimizer

Granddaddy of all existing optimizers

don't go for best plan, go for *least worst plan*

2 Big Ideas

1.

2.



Cost Estimation

`estimate(operator, inputs, stats) → cost`

estimate cost for each operator

depends on input *cardinalities* (# tuples)

discussed earlier in lecture

estimate output size for each operator

need to call `estimate()` on inputs!

use selectivity. assume attributes are independent

Try it in PostgreSQL: `EXPLAIN <query>;`

In order to estimate the cost of an operator, we need to be able to estimate the SIZE of its inputs.

Well the inputs may be other operators, so we need the ability to estimate output sizes

you'll note that I don't have any unit for the cost. That's because estimators tend to be so wildly off that they bear almost no relation to wall clock time.

If you

Estimate Size of Output

```

SELECT      *
FROM        R1, ..., Rn
WHERE       term1 AND ... AND termm
    
```

Query input size

$$|R_1| * \dots * |R_n|$$

Term selectivity

$$\text{col} = v \quad 1 / \text{ICARD}_{\text{col}}$$

$$\text{col1} = \text{col2} \quad 1 / \max(\text{ICARD}_{\text{col1}}, \text{ICARD}_{\text{col2}})$$

$$\text{col} > v \quad (\max_{\text{col}} - v) / (\max_{\text{col}} - \min_{\text{col}})$$

Query output size

$$|R_1| * \dots * |R_n| * \text{term}_1 \text{selectivity} * \dots * \text{term}_m \text{selectivity}$$

ex: suppose emp has 1000 records, dept has 10 records

total records is 1000 * 10, selectivity is 1/1000, so 10 tuples expected to pass join

(note that this is wrong if doing key/fk join on emp.did = dept.did, which will produce 1000 results!)

Note that selectivity is defined relative to size of cross product for joins! p1 and p2

Estimate Size of Output

Emp: 1000 Cardinality

Dept: 10 Cardinality

Cost(Emp join Dept)

Naïve

# total records	$1000 * 10$	$= 10,000$
Selectivity of Emp	$1 / 1000$	$= 0.001$
Selectivity of Dept	$1 / 10$	$= 0.1$
Join Selectivity	$1 / \max(1k, 10)$	$= 0.001$
Output Card:	$10,000 * 0.001$	$= 10$

Key, Foreign Key join

Output Card: 1000

note: selectivity defined wrt cross product size

ex: suppose emp has 1000 records, dept has 10 records

total records is $1000 * 10$, selectivity is $1/1000$, so 10 tuples expected to pass join

(note that this is wrong if doing key/fk join on $\text{emp.did} = \text{dept.did}$, which will produce 1000 results!)

Note that selectivity is defined relative to size of cross product for joins! p1 and p2


Try it out

R.sid = S.sid selectivity 0.01
R.bid selectivity 0.05
 $|R| = M$
 $|S| = N$

Cost: $M + MN$
selection is pipelined

outputs: $0.0005MN$

```
SELECT *  
FROM R, S  
WHERE R.sid = S.sid  
AND R.bid = 10
```



Try it out

R.sid = S.sid selectivity 0.01

R.bid selectivity 0.05


|R| = M

|S| = N

Cost: ?????

outputs: 0.0005MN

```
SELECT *  
FROM R, S  
WHERE R.sid = S.sid  
AND R.bid = 10
```



Try it out

R.sid = S.sid selectivity 0.01

R.bid selectivity 0.05


|R| = M

|S| = N

Cost: $M + (0.05MN)$

outputs: 0.0005MN

```
SELECT *  
FROM R, S  
WHERE R.sid = S.sid  
AND R.bid = 10
```



Selinger Optimizer


Granddaddy of all existing optimizers

don't go for best plan, go for *least worst plan*

- Cost Estimator
 - "predict" cost of query from statistics
 - Includes CPU, disk, memory, etc (can get sophisticated!)
 - It's an art
- Plan Space
 - avoid cross product
 - push selections & projections to leaves as much as possible
 - only join ordering remaining**

Join Plan Space

$A \bowtie B \bowtie C$



How many
plans?

(AB)C (AC)B (BC)A (BA)C (CA)B (CB)A
A(BC) A(CB) B(CA) B(AC) C(AB) C(BA)

parenthetizations * #strings

$\underbrace{\hspace{1.5cm}}_{N!}$

when talking about join orderings, we are ignoring the other operators (e.g., selection, projection), so a query plan is a single binary tree where each node is a join

#strings: For abc, how many distinct strings can these characters create?

#parens: For a given string, how many ways to construct a binary tree on top of that string?

Each column is a possible string ordering

2 possible parenthetizations for three tables

$1 * 2 * 3 = 3!$ String orderings

$6 * 2$ possible join orders for three tables.

What about four?

Join Plan Space

parenthetizations * #strings

A: (A)

AB: (AB)

ABC: ((AB)C), (A(BC))

ABCD: (((AB)C)D), ((A(BC))D), ((AB)(CD)), (A((BC)D)), (A(B(CD)))

paren(n) choose(2(N-1), (N-1)) / N

(choose(2(N-1), (N-1)) / N) * N!

N=10 #plans = 17,643,225,600

The Art of Computer Programming, Volume 4A, page 440-450

How many possible parenthetizations?

$\Rightarrow n! * \text{choose}(2(N-1), (N-1)) / (N)$

$\Rightarrow 4 \text{ choose } 2 / 3 == 6 / 3 = 2$

$6 * 2 == 12$ for 3 relations

$\text{choose}(2*(5-1), (5-1)) / 6$

Selinger Optimizer

Simplify the set of plans so it's tractable and ~ok

1. Push down selections and projections
2. Ignore cross products (S&T don't share attrs)
3. Left deep plans only
4. Dynamic programming optimization problem
5. Consider interesting sort orders

additional simplifying assumptions to make this tractable
At left deep plans only, still $N!$ Possible table orders


This is a classic interview question, one way to solve is to simplify into a dynamic programming problem.

Finally run this for different sort orders that may make future joins or operators cheaper. Eg if there is an order by clause, or downstream join table is already sorted.

Selinger Optimizer

$\text{parens}(N) = 1$

Only left-deep plans
ensures pipelining



Dynamic Programming

Idea: If considering $((ABC)DE)$
compute best (ABC) , cache, and reuse
figure out best way to combine with (DE)

Dynamic Programming Algorithm

compute best join size 1, then size 2, ...
 $\sim O(2^N)$

left deep: first join the two left most tables, then join the result with the next left-most table, and so on.

for a give node, only its left child is allowed to be a join operator

pipelining: right deep tree on the right: before A can join with it's inner relation (result of B join C), need to wait until B join C is completed

With left deep plan, the righ side of join is always a base table that is immediately available.

Note this applies to JOIN ORDERING. Selections and projects can still happen on the right side

Still pretty expensive, since the number of plans is $\text{parens}=1 * N!$

Reducing the Plan Space

Dynamic Programming Algorithm

compute best join size 1, then size 2, ...

R = relations to join

N = |R|

for i in {1, ... N}

for S in {all size i subsets of R}

bestjoin(S) = S-A join A

where A is single relation that minimizes:

cost(bestjoin(S-A)) +

min cost to join A to (S-A) +

min access cost for A

Recursive algorithm

In first iteration, look at all size one subsets of R.

For each table A in the subset, check how it can join with the rest of the subset. For n=1, basically what is the best access path for A

For next iteration, look at all size two subsets. Say PQ, then check:

A=P: cost(bestjoin(Q)) + costofbestjoinoperator(P, Q) + minaccesscost(P).

Bestjoin(Q) already computed

A=Q. Same procedure

Selinger Algorithm $i = 1$

bestjoin(ABC), only nested loops join

$i = 1$

A = best way to access A

B = best way to access B

C = best way to access C

cost: N relations

Selinger Algorithm $i = 2$

bestjoin(ABC)

$i = 2$

$A, B = (A)B \text{ or } (B)A$

$A, C = (A)C \text{ or } (C)A$

$B, C = (B)C \text{ or } (C)B$

cost: $\text{choose}(N, 2) * 2$

There are $\text{choose } n, 2$ ways to pick a set of size 2. Then there are two possible orders of the joins.

Selinger Algorithm $i = 3$

bestjoin(ABC)

$i = 3$

A,B,C = bestjoin(BC)A or
bestjoin(AC)B or
bestjoin(AB)C

cost: choose(N, 3) * 3

Selinger Algorithm Cost

cost = # subsets * # options per subset
 set of relations R
 $N = |R|$

#subsets = choose(N, 1) + choose(N, 2) + choose(N, 3) ...
 $= 2^N$

#options = $k < N$ ways to remove a relation A +
 1 way to join A with R-A (if only NLJ)
 $< N$

Cost = $N * 2^N$
 $N = 12 \quad 49152$

1 way to join A with R-A only because considering left deep NL join only. otherwise it should be M, where M is the number of possible join algorithms

Summary

Single operator optimizations

- Access paths
- Primary vs secondary index costs
- Projection/distinct
- Predicate/project push downs

2 operators aka Joins

- Nested loops, index nested loops, sort merge

Full plan optimizations

- Naïve vs Selinger join ordering

Selectivity estimation

- Statistics and simple models

Summary

Query optimization is a deep, complex topic

Pipelined plan execution

Different types of joins

Cost estimation of single and multiple operators

Join ordering is hard!