

LI0

Transactions, Concurrency, Recovery

Eugene Wu
Fall 2015

Overview

Why do we want transactions?

What guarantees do we want from transactions?

Why Transactions?

Concurrency (for performance)

N clients, no concurrency

1st client runs fast

2nd client waits a bit

3rd client waits a bit longer

Nth client walks away

N clients, concurrency

client 1 runs $x += y$

client 2 runs $x -= y$

what happens?

Can we prevent stepping on toes? *Isolation*

```
x += y
a1 = read(x)
b1 = read(y)
store(a1 + b1)
x -= y
a2 = read(x)
b2 = read(y)
store(a2 - b2)
```

```
x += y
a1 = read(x)
a2 = read(x)
b2 = read(y)
store(a2 - b2)
b1 = read(y)
store(a1 + b1)
```

Why Transactions?

What about 1 client, no concurrency?

Client runs big update query

update set $x += y$

Power goes out

What is the state of the database?

Why Transactions?

What about 1 client, no concurrency?

Client runs big update query

update set $x += y$

Aborts the query (e.g., ctrl-c)

What is the state of the database?

If an abort happens, can the database recover to something sensible? *Atomicity, Durability*

Transactions

Transaction: a sequence of actions

action = read object, write object, commit, abort

API between app semantics and DBMS's view

User's view

T1: begin A=A+100 B=B-100 END

T2: begin A=1.5*A A=1.5*B END

DBMS's logical view

T1: begin r(A) w(A) r(B) w(B) END

T2: begin r(A) w(A) r(B) w(B) END

Transaction Guarantees

Atomicity

users never see in-between xact state.
only see a xact's effects once it's committed

Consistency

database always satisfies ICs.
xacts move from valid database to valid database

Isolation:

from xact's point of view, it's the only xact running

Durability:

if xact commits, its effects *must persist*

Concepts

Concurrency Control

techniques to ensure **correct** results when running transactions concurrently

what does this mean?

Recovery

On crash or abort, how to get back to a consistent (**correct**) state?

The two are intertwined! The CC mechanism dictates the complexity of recovery!

What is Correct?

Serializability

Regardless of the interleaving of operations, end result same as a serial ordering

Schedule

One specific interleaving of the operations

Serial Schedules

Logical xacts

T1: r(A) w(A) r(B) w(B)

T2: r(A) w(A) r(B) w(B)

No concurrency (**serial 1**)

T1: r(A) w(A) r(B) w(B)

T2:

r(A) w(A) r(B) w(B)

No concurrency (**serial 2**)

T1:

r(A) w(A) r(B) w(B)

T2: r(A) w(A) r(B) w(B)

Are serial 1 and serial 2 equivalent?

More Example Schedules

Logical xacts

T1: r(A) w(A) **r(A)** w(B)

T2: r(A) w(A) r(B) w(B)

Concurrency (bad)

T1: r(A) w(A) r(A) w(B)

T2: r(A) w(A) r(B) w(B)

Concurrency (same as serial !!)

T1: r(A) w(A) r(A) w(B)

T2: r(A) w(A) r(B) w(B)

Concepts

Serial schedule

single threaded model. no concurrency.

Equivalent schedule

the database state same at end of both schedules

Serializable schedule (gold standard)

equivalent to a serial schedule

SQL → R/W Operations

```
UPDATE  accounts
SET     bal = bal + 1000
WHERE   bal > 1M
```

Read all balances for every tuple

Update those with balances > 1000

Does the access method matter?

Why Serializable Schedule? Anomalies

Reading in-between (uncommitted) data

T1: R(A) W(A) R(B) W(B) abort

T2: R(A) W(A) commit

WR conflict or dirty reads

Reading same data gets different values

T1: R(A) R(A) W(A) commit

T2: R(A) W(A) commit

RRV conflict or unrepeatable reads

Why Serializable Schedule? Anomalies

Stepping on someone else's writes

T1: W(A) W(B) commit

T2: W(A) W(B) commit

WW conflict or lost writes

Notice: all anomalies involve writing to data that is read/written to.

If we track our writes, maybe can prevent anomalies

Conflict Serializability

What is a conflict?

For 2 operations, if run in different order, get different results

Conflict?	R	W
R	NO	YES
W	YES	YES

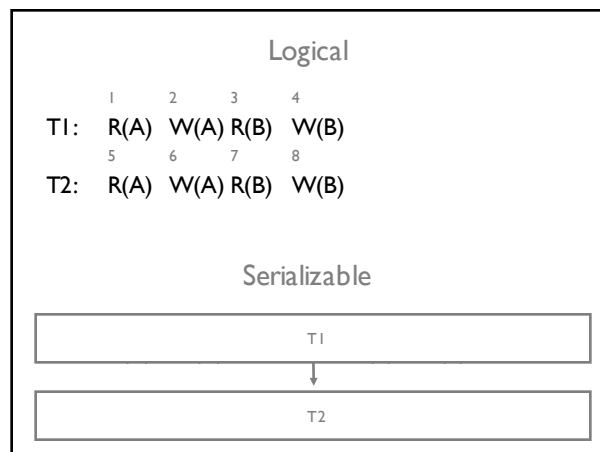

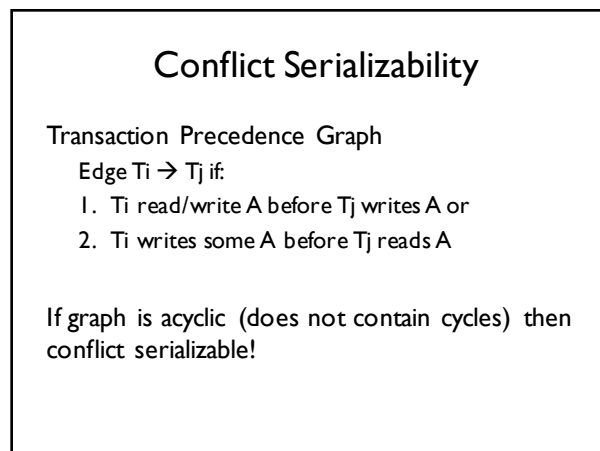
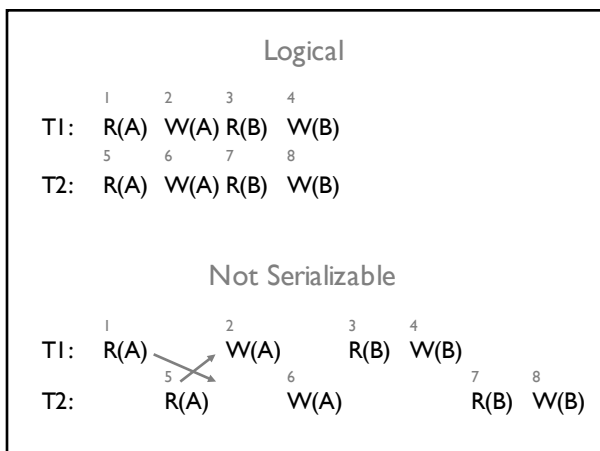
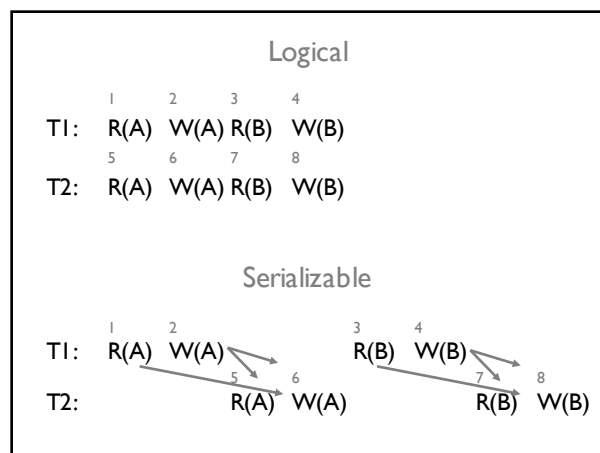
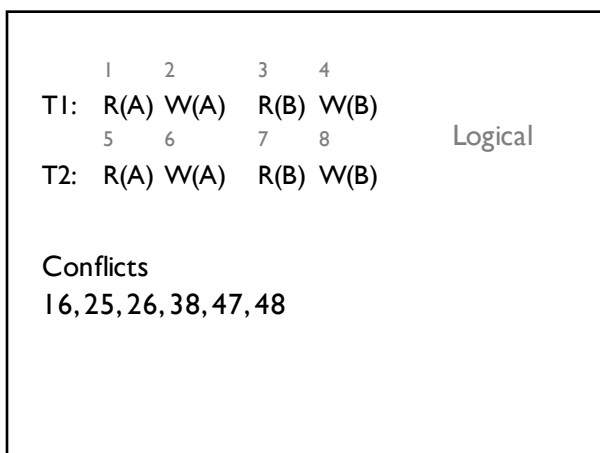
Conflict Serializability


def: possible to swap non-conflicting operations to derive a serial schedule.

∀ conflicting operations O1 of T1, O2 of T2

O1 always before O2 in the schedule or

O2 always before O1 in the schedule





Fine, but what about COMMITing?

T1 R(A) W(A) R(B) ABORT
T2 R(A) COMMIT

Not recoverable

Promised T2 everything is OK. IT WAS A LIE.

T1 R(A) W(B) W(A) ABORT
T2 R(A) W(A)

Cascading Rollback.

T2 read uncommitted data → T1's abort undoes T1's ops & T2's

Lock-based Concurrency Control

Must get a shared(read) or exclusive(write) lock BEFORE op
If other xact has lock, can get if lock table says so

YES

			T1
	Allowed?	S	X
T2		S	Y
		X	N
		X	N

Can this schedule happen?

T1 R(A) W(A) R(B) ABORT
T2 R(A) COMMIT

Lock-based Concurrency Control

Two-phase locking (2PL)

Growing phase: acquire locks

Shrinking phase: release locks

T1 R(A) W(B) W(A) ABORT
T2 R(A) W(A)

shrink here

Uh Oh, same problem

Lock-based Concurrency Control

Strict two-phase locking (Strict 2PL)

Growing phase: acquire locks

Shrinking phase: release locks

Hold onto locks until commit/abort



Why? Which problem does it prevent?

T1 R(A) W(B) W(A) ABORT
T2 R(A) W(A)

Guarantees serializable schedules! Avoids cascading rollbacks!

Review

Issues

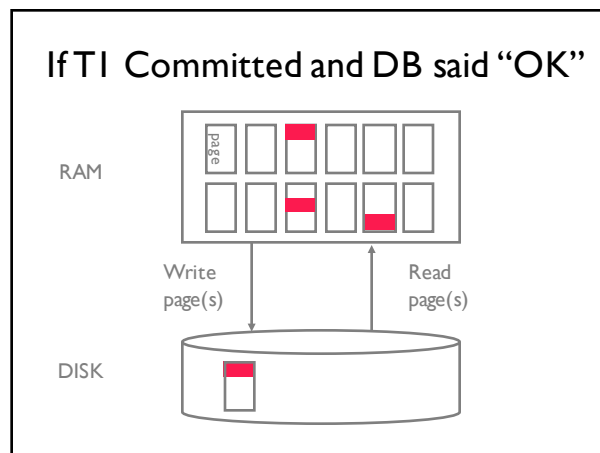



TR: dirty reads
RW: unrepeatable reads
WW: lost writes

Schedules

Equivalence
Serial
Serializable

Serializability

Conflict serializability
how to detect
Conflict Serializable Issues
Not recoverable
Cascading Rollback
Strict 2 phase locking



Recovery

Two properties: Atomicity, Durability

Assumption in class: disk is safe. memory is not.

Need to account for
 when pages are modified
 when pages are flushed to disk

Recovery

Deal with 2 cases

If T2 commits, what could make it not durable?
 didn't write all changed pages to disk

When could uncommitted ops appear after crash?

wrote modified pages before commit

Aborts and Undos

If Tx aborts, all of its actions must be undone.

Ty that read Tx's writes must be aborted
(cascading abort)

Strict 2PL avoids cascading aborts

Use a log to know what actions to undo

```
1. A = 1
2. B = 5
3. C = 10
4. BEGIN T5
5. A = 10
6. B = B + A
7. C = B - 2
8. ABORT
9. undo 7
10. undo 6
...
```

Aborts and Undos

If Tx aborts, all of its actions must be undone.

Ty that read Tx's writes must be aborted
(cascading abort)

Strict 2PL avoids cascading aborts

Use a log to know what actions to undo

On crash, abort all non-committed xacts

```
1. A = 1
2. B = 5
3. C = 10
4. BEGIN T5
5. A = 10
6. B = B + A
7. CRASH
```

Logs

Log records

writes: old & new value

commit/abort actions

xact id & xact's previous log record

Write ahead logging (WAL)

log records stored on disk (persisted) *before* data pages can be persisted

log is the *ground truth*

So far: use log to undo partial transactions

Durability

Bad scenario

T1 writes to A

T1 commits, log record written to disk

start writing page with A to disk

crash

Can undo help us?

Need to redo T1, otherwise no durability!

Aries Recovery Algorithm

3 phases

1. Analyze the log to find status of all xacts
2. Redo xacts that were committed
3. Undo partial xacts

Recovery is *extremely* tricky and *must be correct*