L10
Transactions, Concurrency, Recovery

Eugene Wu
Fall 2015

Overview

Why do we want transactions?

What guarantees do we want from transactions?

Why Transactions?

Concurrency (for performance)
- N clients, no concurrency
  - 1st client runs fast
  - 2nd client waits a bit
  - 3rd client waits a bit longer
  - Nth client walks away
- N clients, concurrency
  - client 1 runs x += y
  - client 2 runs x -= y
  - what happens?

Can we prevent stepping on toes? *Isolation*

x += y
a1 = read(x)
b1 = read(y)
store(a1 + b1)
x -= y
a2 = read(x)
b2 = read(y)
store(a2 – b2)

x += y
a1 = read(x)
a2 = read(x)
b2 = read(y)
store(a2 – b2)
b1 = read(y)
store(a1 + b1)

Why Transactions?

What about 1 client, no concurrency?
- Client runs big update query
  - update set x += y
- Power goes out
- What is the state of the database?

Why Transactions?

What about 1 client, no concurrency?
- Client runs big update query
  - update set x += y
- Aborts the query (e.g., ctrl-c)
- What is the state of the database?

If an abort happens, can the database recover to something sensible? *Atomicity, Durability*

Transactions

Transaction: a sequence of actions
- action = read object, write object, commit, abort
- API between app semantics and DBMS's view

User's view

T1: begin	A=A+100	B=B-100	END
T2: begin	A=1.5*A	A=1.5*B	END

DBMS's logical view

T1: begin	r(A) w(A)	r(B) w(B)	END
T2: begin	r(A) w(A)	r(B) w(A)	END

Transaction Guarantees

Atomicity
- users never see in-between xact state.
- only see a xact's effects once it's commited

Consistency
- database always satisfies ICs.
- xacts move from valid database to valid database

Isolation:
- from xact's point of view, it's the only xact running

Durability:
- if xact commits, its effects *must persist*

Concepts

Concurrency Control
- techniques to ensure correct results when running transactions concurrently

what does this mean?

Recovery
- On crash or abort, how to get back to a consistent (correct) state?

The two are intertwined! The CC mechanism dictates the complexity of recovery!

What is Correct?

Serializability
- Regardless of the interleaving of operations, end result same as a serial ordering

Schedule
- One specific interleaving of the operations

Serial Schedules

Logical xacts

T1: r(A) w(A) r(B) w(B)
T2: r(A) w(A) r(B) w(B)

No concurrency (serial 1)

T1:	r(A) w(A)	r(B) w(B)		
T2:			r(A) w(A)	r(B) w(B)

No concurrency (serial 2)

T1:			r(A) w(A)	r(B) w(B)
T2:	r(A) w(A)	r(B) w(B)		

Are serial 1 and serial 2 equivalent?

More Example Schedules

Logical xacts

T1: r(A) w(A) r(A) w(B)
T2: r(A) w(A) r(B) w(B)

Concurrency (bad)

T1:	r(A) w(A)		r(A) w(B)	
T2:		r(A) w(A)		r(B) w(B)

Concurrency (same as serial 1!)

T1:	r(A) w(A)		r(A) w(B)	
T2:		r(A)		w(A) r(B) w(B)

Concepts

Serial schedule

single threaded model. no concurrency.

Equivalent schedule

the database state same at end of both schedules

Serializable schedule (gold standard)

equivalent to a serial schedule

SQL → R/W Operations

```
UPDATE    accounts
SET       bal  =  bal  +  1000
WHERE     bal  >  1M
```

Read all balances for every tuple

Update those with balances > 1000

Does the access method mater?

Why Serializable Schedule? Anomalies

Reading in-between (uncommitted) data

T1: R(A) W(A) R(B) W(B) abort

T2: R(A) W(A) commit

WR conflict or dirty reads

Reading same data gets different values

T1: R(A) R(A) W(A) commit

T2: R(A) W(A) commit

RW conflict or unrepeatable reads

Why Serializable Schedule? Anomalies

Stepping on someone else's writes

T1: W(A) W(B) commit

T2: W(A) W(B) commit

WW conflict or lost writes

Notice: all anomalies involve writing to data that is read/written to.

If we track our writes, maybe can prevent anomalies

Conflict Serializability

What is a conflict?

For 2 operations, if run in different order, get different results

Conflict?	R	W
R	NO	YES
W	YES	YES

Conflict Serializability

def: possible to swap non-conflicting operations to derive a serial schedule.

∀ conflicting operations O1 of T1, O2 of T2

O1 always before O2 in the schedule or

O2 always before O1 in the schedule

	1	2	3	4
T1:	R(A)	W(A)	R(B)	W(B)
	5	6	7	8
T2:	R(A)	W(A)	R(B)	W(B)

Logical

Conflicts
16, 25, 26, 38, 47, 48

Logical

	1	2	3	4
T1:	R(A)	W(A)	R(B)	W(B)
	5	6	7	8
T2:	R(A)	W(A)	R(B)	W(B)

Serializable

	1	2			3	4		
T1:	R(A)	W(A)			R(B)	W(B)		
			5	6			7	8
T2:			R(A)	W(A)			R(B)	W(B)

Logical

	1	2	3	4
T1:	R(A)	W(A)	R(B)	W(B)
	5	6	7	8
T2:	R(A)	W(A)	R(B)	W(B)

Not Serializable

	1		2		3	4		
T1:	R(A)		W(A)		R(B)	W(B)		
		5		6			7	8
T2:		R(A)		W(A)			R(B)	W(B)

Conflict Serializability

Transaction Precedence Graph

Edge Ti → Tj if:

1. Ti read/write A before Tj writes A or
2. Ti writes some A before Tj reads A

If graph is acyclic (does not contain cycles) then conflict serializable!

Logical

	1	2	3	4
T1:	R(A)	W(A)	R(B)	W(B)
	5	6	7	8
T2:	R(A)	W(A)	R(B)	W(B)

Serializable

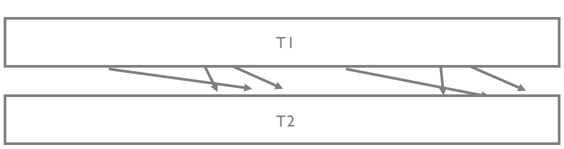

Logical

	1	2	3	4
T1:	R(A)	W(A)	R(B)	W(B)
	5	6	7	8
T2:	R(A)	W(A)	R(B)	W(B)

Serializable

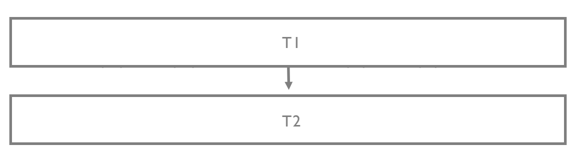

Logical

	1	2	3	4
T1:	R(A)	W(A)	R(B)	W(B)
	5	6	7	8
T2:	R(A)	W(A)	R(B)	W(B)

Not Serializable

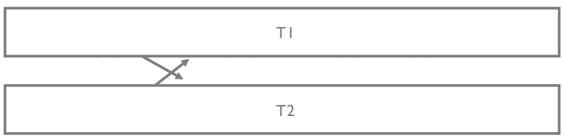

Fine, but what about COMMITing?

T1 R(A) W(A) R(B) ABORT
T2 R(A) COMMIT

Not recoverable

Promised T2 everything is OK. IT WAS A LIE.

T1 R(A) W(B) W(A) ABORT
T2 R(A) W(A)

Cascading Rollback.

T2 read uncommitted data → T1's abort undos T1's ops & T2's

Lock-based Concurrency Control

Must get a shared(read) or exclusive(write) lock BEFORE op
If other xact has lock, can get if lock table says so

YES

		T1	
	Allowed?	S	X
T2	S	Y	N
	X	N	N

Can this schedule happen?

T1 R(A) W(A) R(B) ABORT
T2 R(A) COMMIT

Lock-based Concurrency Control

Two-phase locking (2PL)
- Growing phase: acquire locks
- Shrinking phase: release locks

shrink here

T1 R(A) W(B)W(A) ABORT
T2 R(A) W(A)

Uh Oh, same problem

Lock-based Concurrency Control

Strict two-phase locking (Strict 2PL)
- Growing phase: acquire locks
- Shrinking phase: release locks
- Hold onto locks until commit/abort

Why? Which problem does it prevent?

T1 R(A) W(B) W(A) ABORT
T2 R(A) W(A)

Guarantees serializable schedules! Avoids cascading rollbacks!

Review

Issues
- TR: dirty reads
- RW: unrepeatable reads
- WW: lost writes

Schedules
- Equivalence
- Serial
- Serializable

Serializability
- Conflict serializability
- how to detect

Conflict Serializable Issues
- Not recoverable
- Cascading Rollback

Strict 2 phase locking

Normal Execution

CRASH

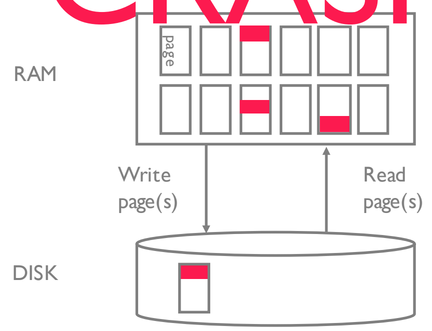

After a Crash

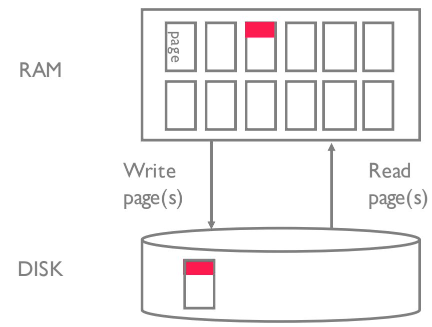

If DB did not say "OK, committed"

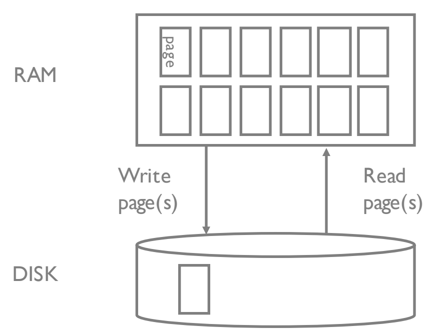

If T1 Committed and DB said "OK"

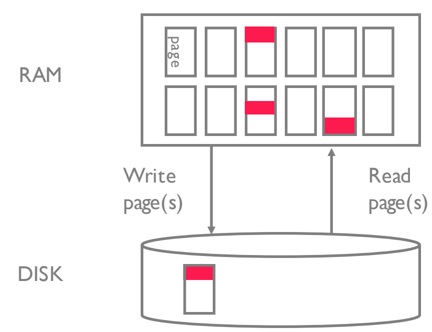

Recovery

Two properties: Atomicity, Durability

Assumption in class
- Disk is safe. Memory is not.
- Running strict-2PL

Need to account for
- when pages are modified
- when pages are flushed to disk

Recovery

Deal with 2 cases

If T2 commits, what could make it not durable?
- didn't write all changed pages to disk

When could uncommitted ops appear after crash?
- wrote modified pages before commit

Aborts and Undos

If Tx aborts, must undo all its actions
- Ty that read Tx's writes must be aborted (cascading abort)
- Strict 2PL avoids cascading aborts

Use a log to know what actions to undo

1. A = 1
2. B = 5
3. C = 10
4. BEGIN T5
5. A = 10
6. B = B + A
7. C = B – 2
8. ABORT
9. undo 7
10. undo 6

...

Aborts and Undos

If Tx aborts, must undo all its actions
- Ty that read Tx's writes must be aborted (cascading abort)
- Strict 2PL avoids cascading aborts

Use a log to know what actions to undo
On crash, abort all non-committed xacts

1. A = 1
2. B = 5
3. C = 10
4. BEGIN T5
5. A = 10
6. B = B + A
7. CRASH

Logs

Log is the *ground truth*

Log records
- writes: old & new value
- commit/abort actions
- xact id & xact's previous log record

Persist log records (write to disk) *before* data pages persisted
Is this enough?

Durability

Bad scenario
- T1 writes to A
- T1 commits, log record written to disk
- start writing page with A to disk
- *crash*

Can undo help us?
Need to redo T1, otherwise no durability!

Logs

Log is the *ground truth*

Log records
- writes: old & new value
- commit/abort actions
- xact id & xact's previous log record

Write ahead logging (WAL)
1. Persist log records (write to disk) *before* data pages persisted
2. Persist all log records *before* commit

(1) guarantees UNDO info
(2) guarantees REDO info

Aries Recovery Algorithm

3 phases
- Analyze the log to find status of all xacts
  - Commited or in flight?
- Redo xacts that were committed
- Undo partial (in flight) xacts

Recovery is *extremely* tricky and *must be correct*

Aries

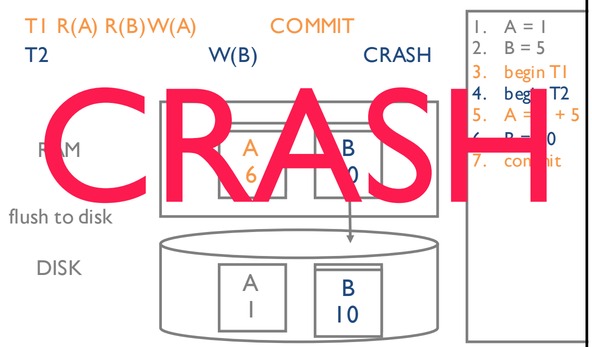

Aries: alternative flushing order

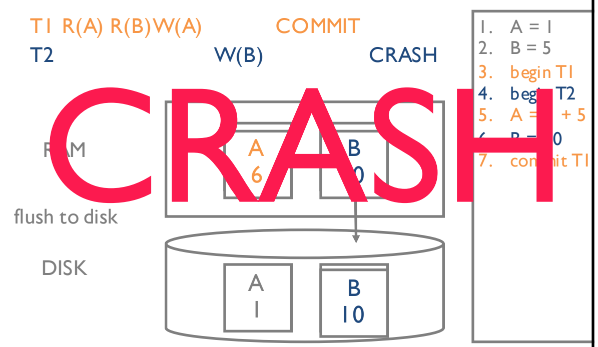

Aborts and Undos

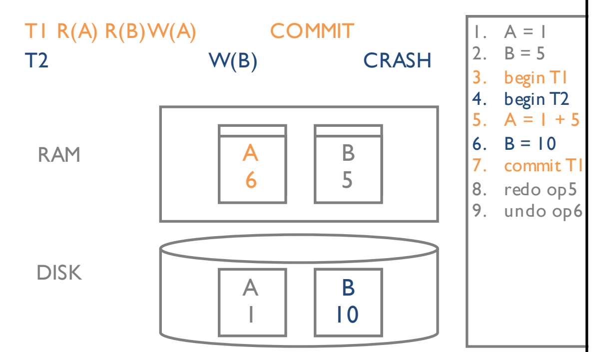

Summary

Recovery depends on what failures are tolerable

Buffer pool can write RAM pages to disk any time

Recovery Manager ensures durability and atomicity via redo and undo

You should know

What transactions/schedules/serializable are
Can identify conflict serializable schedules
Can identify schedule anomalies
Can identify strict 2PL executions

Understand WAL and what it provides
Given an executed schedule, run the proper sequence of undo/redos